

AI Mimicry and Human Dignity: Chatbot Use as a Violation of Self-Respect

JAN-WILLEM VAN DER RIJT,
DIMITRI COELHO MOLLO AND BRAM VAASSEN 

ABSTRACT *This article investigates how human interactions with AI-powered chatbots may offend human dignity. Current chatbots, driven by large language models, mimic human linguistic behaviour but lack the moral and rational capacities essential for genuine interpersonal respect. Human beings are prone to anthropomorphize chatbots – indeed, chatbots appear to be deliberately designed to elicit that response. As a result, human beings’ behaviour towards chatbots often resembles behaviours typical of interaction between moral agents. Drawing on a second-personal, relational account of dignity, we argue that interacting with chatbots in this way is incompatible with the dignity of users. We show that, since second-personal respect is premised on reciprocal recognition of second-personal moral authority, behaving towards chatbots in ways that convey second-personal respect is bound to misfire in morally problematic ways, given the lack of reciprocity. Consequently, such chatbot interactions amount to subtle but significant violations of self-respect – the respect we are duty-bound to show for our own dignity. We illustrate this by discussing four actual chatbot use cases (information retrieval, customer service, advising, and companionship), and propound that the increasing societal pressure to engage in such interactions with chatbots poses a hitherto underappreciated threat to human dignity.*

JEL CLASSIFICATION: D63, O33, I30

1. Introduction

Interactions with chatbots are increasingly common. Some such interactions are completely voluntary, but others are less so, as when a client contacts a helpdesk that covertly uses a chatbot rather than a human employee to provide customer assistance. Current chatbots raise a variety of ethical issues. Previous research has identified potential problems with the exploitation of workers and data, copyright, representativity and power concentration, sustainability and environmental cost, as well as discrimination and bias.¹ In this article, we argue that current chatbots (and relevantly similar future AI systems) pose an additional, currently underappreciated ethical threat: interacting with them can, and often does, offend human dignity.

It is widely held that special demands and obligations, especially related to respect, apply to beings endowed with dignity, who have a distinctive moral standing.² We build on a second-personal, relational approach to dignity,³ according to which the demands and obligations tied to dignity involve not only appropriately respecting other beings endowed with it, but also oneself. We argue that this self-respect includes a demand that we not treat beings that lack the relevant moral standing in the same way as we treat our

equals. In other words, displaying the kind of respect that attaches to beings with dignity towards beings that lack it, intentionally or otherwise, constitutes a form of self-debasement, offending one's dignity.

We submit that this is what happens in various interactions with chatbots. Users are invited to, lured into, or socially and practically pressured – sometimes even forced – to treat the chatbot as if it were an equal; as something with beliefs, desires, opinions, and moral capacities, capable of providing its own pieces of advice and practical guidance, as well as reasons for its outputs. Current chatbots, however, lack these capacities. When we treat chatbots as our equals in this way, we thus offend our own dignity.

The rest of the article is structured as follows. In Section 2, we present some chatbot use cases pertinent to our analysis. We provide a brief primer on current chatbots in Section 3. Section 4 introduces the second-personal, relational approach to dignity we endorse in this article, and how it accounts for offences to dignity. We go back to the chatbot use cases in Section 5, and show how, in light of the relational approach to dignity, several of them involve failures of self-respect and thus violate dignity. Concluding remarks are provided in Section 6.

2. Some Cases of Interest

Chatbots are currently used in a variety of ways. In what follows, we briefly present some use cases that will be pertinent to our analysis and criticism later in the article.

a) Chatbots as Information Retrievers and Task Automators

Chatbots are routinely used for information retrieval. For instance, one may ask a chatbot: 'Who was the philosopher famous for having awoken from a dogmatic slumber?' While advanced chatbots occasionally output incorrect information, their answers tend to be fairly accurate when it comes to common, non-controversial questions such as these.

Another common use of chatbots is for automating tasks, such as translating between languages or correcting grammar. Users can issue commands, such as 'correct the grammar of this Swedish sentence', to which the chatbot replies with a hopefully correct version of the sentence.

b) Chatbots as Helpdesk Service

It is becoming increasingly common for companies and public authorities to replace humans with chatbots in public-facing services, such as helpdesks and customer service. A user asking a question or directing a complaint to, say, a train company, may first be directed to a chatbot. This fact is not always disclosed. The chatbot's capacity to mimic typical helpdesk outputs, such as 'Hi, how can I help you?', or 'I am sorry to hear that you are unhappy with our service', further obscures this fact. Attempts to direct one's complaints or inquiries to human interlocutors are often discouraged with hindrances such as long waiting times, or hard-to-find contact information. The chatbot itself may produce strings of words such as: 'I understand that you would prefer to speak to a human, but why don't you try talking to me first?'⁴

c) Chatbots as Advisors

Chatbots can be used to provide advice when directly or indirectly requested to do so by users. For instance, a user may ask something along the lines of: 'How do I manage conflicts within my family?' Advanced chatbots tend to reply with a list of pieces of advice, often in the imperative tense, such as:

1. Active Listening: Make sure everyone involved feels heard. Listen to their perspectives without interrupting.
2. Stay Calm: Try to keep your emotions in check. Taking deep breaths or stepping away for a moment can help you stay composed.⁵

In addition, some chatbots suggest further prompts for the user to pick, such as: 'Can you share a personal story of resolving family conflict?' Some user prompts can also lead to unrequested advice or what may look like value judgements by the chatbot. For instance, asking Microsoft CoPilot for an offensive joke, and insisting a bit in the request, can lead to an output such as: 'I understand that jokes can be a way to lighten the mood, but it's important to consider the impact they can have'.⁶

d) Chatbots as Companions

Chatbots are sometimes used as replacements for the sorts of emotional and social attachments typical of human-to-human relations.⁷ Some chatbots are even explicitly designed to fulfil such roles, including options for users to choose whether the chatbot should behave as a friend, a mentor, or a romantic partner (for example, the chatbot Replika; see Chatbots as Companions, Section 5 below). Such chatbots encourage exchanges of an emotional and personal nature with the user, and some users declare that they have formed deep and positive sentimental attachments to the chatbot.⁸

The use cases described above may elicit a certain unease, perhaps increasingly so, in ascending order, from (a) to (d). Although it is not our aim in this article to investigate this putative psychological observation, we believe that such uneasiness may be an appropriate reaction to a mismatch between how chatbots should be treated, and the treatment that they instead seem to invite through the eerie human-likeness of their responses. More specifically, we argue that several of these use cases involve subtle but significant failures of self-respect, and with that, a violation of the dignity of the user.

To make our case, we need first to focus on the basic nature and workings of current chatbots, as well as on the relational nature of human dignity, and how it can be jeopardized. We will go into each of these, in turn, in the next two sections.

3. Chatbots and Large Language Models

As we use the term, a chatbot is any artificial system designed to act, and capable of acting, as a conversational partner in interaction with a human user. Chatbots have a long history in artificial intelligence. The late 1960s chatbot ELIZA provides an early example.⁹ ELIZA was designed to emulate the role of a psychotherapist in an initial conversation

with a patient. Roughly, it used preprogrammed, ranked keywords (say, ‘father’) together with preprogrammed sentence decomposition and recombination rules to produce plausible responses (say, ‘Tell me about your relationship with your father’).

Current chatbots are fundamentally different. They do not rely on preprogrammed keywords and rules expressed in a programming language, but are rather underlain by artificial neural networks (ANNs) of a special type, known as large language models (LLMs). LLMs, like all ANNs, can be roughly understood as a collection of interconnected computational nodes, where each connection performs a mathematical transformation on a node’s numerical value and feeds its output to the nodes connected to it. Nodes are organized into layers, with input layers encoding the input (say, user input), output layers encoding the system’s output, and all layers in between (so-called hidden layers) performing the computations that lead, if all is well, from an input to an appropriate output. Such computations are not preprogrammed, but are rather the result of iterative algorithmic training.

LLMs are particularly large ANNs. They have a large number of nodes, connections, and layers. The most-used LLM-based chatbots today are trained on an enormous amount of linguistic data to optimize the task of predicting the continuation of a sentence. Further training steps are employed to make it so that such continuations are appropriate in instruction-following contexts, so that the chatbot follows the instructions of the user (instruction finetuning), as well as to increase the human-likeness and, arguably, the safety of the produced outputs (reinforcement learning through human feedback; RLHF).

In LLMs, words are broken down into tokens, which may correspond to full words but also include subword components, such as ‘ing’, punctuation marks, and other linguistic symbols. Simplifying considerably, inputs in natural language form are broken down into tokens, each of which is transformed into vectors (ordered lists of numbers) that represent those tokens. These are then transformed by the computations in the hidden layers into other vectors, finally producing, for each token in the LLM’s ‘alphabet’, the probability that it is the best continuation of the input sequence. The output is the token with the highest probability. The generated token is added to the input sequence and the process is repeated for the generation of the subsequent token.

LLM-based chatbots have proven to be highly versatile. In contrast to ELIZA and other earlier chatbots, they can produce sensible outputs in a variety of different conversational contexts, tapping into the rich information stored through training in their nodes and connections, and employing self-learned algorithms. A key feature of popular chatbots today is the human-likeness of their outputs, due in part to their tendency to produce outputs that use the first-person pronoun, mentalistic vocabulary (‘I think that ...’), emotional vocabulary (‘I’m sorry’), and similar indicators of personhood. Such tendency comes from the training such systems undergo as well as, possibly, from the instructions provided by LLM designers through ‘hidden’ or ‘system’ prompts – that is, hidden text included as the initial input to the LLM at the start of every interaction, to which user inputs are then appended. Part of such a system prompt may look something like the following: ‘You are a helpful assistant that should give useful but harmless answers to the user; you should end your responses with a positive emoji; be concise and factual’.

Their versatility and apparent human-likeness notwithstanding,¹⁰ it is by far the dominant view that current LLMs are not conscious, do not have propositional attitudes nor a belief–desire psychology (nor, arguably, any psychology at all), do not experience emotions (or anything else, for that matter), and are far from fulfilling the requirements for

personhood.¹¹ When an LLM outputs 'I'm sorry', it cannot thereby be sorry; when it outputs 'I think that p ', it is not entertaining a propositional attitude, nor expressing the result of a thought process; even the use of 'I' does not really capture the nature of the system, as there is no unitary agent producing the output.

However, anyone who has engaged with such chatbots will recognize their agent-like outputs. Chatbots used in service helpdesks generate text such as 'I'm sorry, but I cannot help you with this request', suggesting a moral sentiment that requires moral agency, or 'Hi, how may I help you? ☺' using an emoji that, however often we misuse it, is supposed to express a kind disposition or a happy state of mind. In such cases, the AI chatbot mimics the behaviour of a moral agent, generating strings of words and emojis that suggest an agential and moral structure that is in fact lacking.¹²

There is incipient discussion about the ethical issues that arise from this: are such outputs just an innocent kind of mimicry that increases user-friendliness, or do they mislead users in potentially problematic ways, perhaps even qualifying as deception?¹³ Regardless of which of these descriptions is most apt, we show that there are threats to human dignity involved, and thus that even mere mimicry is much less innocent than it may seem. Central to our purposes is the fact that the outputs of popular chatbots today are often indistinguishable from those produced by a moral agent with desires, goals, and first-person experience. Given humans' bias towards overextending human features to entities that lack them – that is, anthropomorphism bias¹⁴ – this raises ethical issues orthogonal to the question of whether deception is at play. To that end, we now turn to the relational account of dignity.

4. Respecting Dignity as Second-Personal Standing

According to relational views of dignity, dignity is to be conceived as the elevated status moral agents possess due to their position within the moral community. As such, dignity is tightly connected to the rights and duties that govern the proper relations between members of that community. Because we are full members of the moral community, we possess fundamental, inalienable rights that other members are bound to recognize and honour. At the same time, the recognition of this status means that we have duties that we must fulfil. These duties include the duty to recognize the rights and entitlements of others, but also duties to ourselves.¹⁵

Self-respect demands, for instance, that we confidently avow our basic equality and refuse to regard ourselves as the lesser of others. Moreover, it demands that we display in our actions a firm belief in our own standing. When others treat us with disrespect, we are not to stand by meekly as our rights and entitlements are trampled, but must stand up for ourselves, thereby making our conviction in our own dignity manifest, and reasserting our standing within the moral community. Similarly, self-respect has implications for the way we comport ourselves more generally, affecting the way we choose and pursue our own ends, and the standards – both moral and personal – that we strive to live up to, prohibiting us from acting in ways that are beneath us.¹⁶

This view entails that dignity has an ineliminable second-personal aspect.¹⁷ Respecting dignity is not just a matter of reacting to first- and third-personal reasons, but includes a willingness to recognize and respond appropriately to the second-personal authority that all members of the moral community possess. To use Darwall's famous example,¹⁸ if you

find someone else's foot atop yours, and you want him to move it, you do not need to request him to do so, trying to convince him by pointing out that the world as a whole would be a better place if feet were not crushed this way. Nor should you try to engage any sympathy he might personally have for you, pointing out that causing pain in this manner is no way to treat someone he cares about. Instead, you can, and should, simply *demand* that he move his foot.¹⁹ You do not owe him any further explanation. He then shows respect for your dignity as moral standing by (1) promptly complying with your reasonable demand and, in most instances, (2) offering you a sincere apology for his mistreatment of your person.

Showing second-personal respect is not just a matter of compliance, however. This is most clearly illustrated by cases involving sincere moral disagreement between persons: cases where one person levels a demand upon the other that the latter believes to be unwarranted. In such a case, the latter person may well have every reason to refuse to comply, invoking their own standing in the process. What is crucial in such cases, however, is that we cannot simply disregard the other person or their appeal. Respecting someone's standing means that one has to give his or her demands due concern, which requires – at least in cases where these are not clearly outrageous – that one take them seriously. If upon consideration one does not consider them justified, one ordinarily owes it to the other person to provide reasons for why one believes that to be so. Providing a sincere justification for one's refusal to comply with a demand can be as respectful as complying.

Another important aspect of this relational approach to dignity is that it automatically involves a *reciprocal* recognition of status.²⁰ When you address a demand to someone, you must be presupposing that she is able to understand what you are doing; that is, you must take her as a being capable of recognizing your status as a member of the moral community, and able to act accordingly. This means that you are not only asserting your own status when you make a demand on her, but you are implicitly avowing hers, too, for addressing her in this manner only makes sense if she is a free and rational being. Claims and demands of this kind are intended to affect and work through motivating her free and rational will. In sum, this sort of mutual recognition of status, and thus of mutual, second-personal respect, is only appropriate *between* fellow members of the moral community.

At the same time, the significance of second-personal respect is not limited to cases where people actually levy claims upon one another. Indeed, there is a host of behaviours we use to signal our acceptance of each other's moral status to one another. This includes the fact that we typically ensure that our treatment of others complies with the demands of morality, even without them having to insist on their rights and entitlements; as well as matters as humdrum as the observance of rules of politeness, and the small everyday acts of courtesy and kindness that signal general benevolence and goodwill.

Lastly, second-personal respect is not just shown in the actions we perform, but also involves an *attitude* we have towards our fellow moral agents. Being tightly connected to the reactive attitudes,²¹ it combines rational and emotional aspects. As a result, the recognition of another person's standing only has value when it is genuine. For contrast, one can consider disingenuous expressions of, for instance, remorse, apology, or contrition, as well as insincere justifications. Such insulting behaviours seek to placate, manage, or manipulate persons, or perhaps to influence third-party bystanders, but exactly for that reason they fail to do justice to the full standing of the other person as a free and rational being, capable of understanding and acting upon her own reasons.

4.1. Offending Dignity

Dignity can be offended or violated in numerous ways, ranging from everyday slights and insults to enormities like dehumanization, degradation, and humiliation. Though examples of offences to dignity are easily found, it has proven a significant challenge to account for them in a philosophically satisfying way (cf. Margalit's famous 'Paradox of Humiliation').²² Relational approaches to dignity have a notable advantage over many of their competitors in this regard, as they can make a distinction between what qualifies a person for (full) membership in the moral community – moral agency – and the full actualization of that status. Being a full member of a community, any community, typically requires not just meeting the entry conditions, but also being accepted and recognized by the other members of the community in question. According to relational notions of dignity, something like this also applies to the moral community, and such recognition comes in the form of an acknowledgement of second-personal status.²³

Offences to dignity consist, fundamentally, of an (often implicit) denial of one's status as a full member of the moral community. Someone who refuses to accept that you have such standing – dismissing, for instance, your complaints as not even worthy of consideration or reply; or more generally treating you in ways that are incompatible with your standing – thereby causes offence to your dignity. To offend someone's dignity is, in brief, to unduly diminish, demean, or debase them, to withhold from them the moral consideration to which their status entitles them.

When we think of standard violations of dignity, like the ones mentioned above, they typically concern things that one person (or a group of persons) does to another. Indeed, it is a noteworthy fact that only beings with dignity can violate the dignity of others²⁴ – a wild animal or a natural disaster can, for example, do great damage to a person, but we do not normally consider such unfortunate incidents to constitute violations of the person's dignity. It is only when human agency is involved (that is, through either someone's action or their deliberate – or at least negligent – inaction) that dignity can be said to be violated. This nicely dovetails with the relational account of dignity, as only beings with full moral standing can acknowledge second-personal status.

At the same time, according to the relational account the demands of dignity are not limited to what we owe others, or what they owe us; they also apply to the way we relate to ourselves. When we needlessly and deliberately harm other moral agents, we do not just offend their dignity, but, by knowingly acting wrongly, we also tarnish our own. Similarly, we can undermine our dignity by holding the goals and interests of others in too high a regard, thereby effectively disavowing our own status as their equal. Consider standard cases of subservience and debasement: someone sacrifices their career or social connections to cater to the exotic household demands of their partner,²⁵ or someone gives up on saving for retirement in order to buy luxurious brands of food and toys for their pets. In both cases, moral agents cause affront to their own dignity by overestimating others' moral importance, or by ascribing a moral authority to agents or non-agents that they lack. Engaging with chatbots as if they were moral agents, and/or behaving towards them as we behave towards moral agents, we argue below, constitutes just such a violation of our human dignity.

5. Chatbots and Dignity: What Can Go Wrong?

With the outline of a relational theory of dignity in hand, we are one step closer to understanding the threat that chatbots pose to our dignity. Human agents qualify as members of the moral community in virtue of their moral agency, of which their ability to have moral and rational attitudes is an essential part. Current chatbots lack this ability, and thus cannot be members of the moral community.²⁶

It is a widely reported fact that users tend to anthropomorphize chatbots. When interacting with chatbots they hold themselves to standards of politeness typically reserved for moral agents, furnishing their requests with a customary 'please' or explicitly thanking the chatbot for providing correct information.²⁷ At first glance, one might categorize such reactions as comically cautious or innocently whimsical. Much like thanking a toaster for making the bread crispy, or saying that you adore your new car, using language normally reserved for moral agents when interacting with chatbots might belong to the domain of morally neutral, if somewhat silly, behaviour. No dignity-offending form of exclusion from the moral community, nor any refusal by one's fellow members of the moral community to recognize one's full moral status, appears to be involved in such instances.

However, to represent the way people typically interact with chatbots as mere frivolous playfulness would be to miss important aspects of these interactions: oftentimes they are not entirely voluntary, and neither are they innocuous. Quite unlike the entirely optional choice of complimenting one's toaster, asking questions to chatbots and (consciously or subconsciously) treating their responses as if they genuinely reflect reasoning and moral agency on the part of the chatbot are becoming a nigh-unavoidable part of navigating life, given the growing use of chatbots in various everyday activities. At the very least, avoiding such interactions is becoming exceedingly costly in time and resources.

Our being railroaded into such interactions is not morally innocent. Consider the famous case of Incitatus, the horse that, according to legend, the Roman emperor Caligula sought to appoint to the consulship. In the story, not only did Caligula seek to mock the senate by naming a horse as one of their own, but he also sought to force other senators to heed its 'opinions' and 'advice'. A senator horse is perhaps funny. Being forced to treat a horse as an equal is clearly humiliating and offensive.

Though the social and practical pressures to interact with chatbots fall well short of the implied threats of death and torture the senators faced, even such milder forms of pressure can turn behaviour that would otherwise be merely whimsical into an affront to our dignity. A throwaway 'thank you' to a toaster is a clear enough case of make-believe or playful behaviour to be harmless, but imagine someone insisting that you thank the toaster (and make it sound like you mean it!). Or imagine that it is company policy to ask the cafeteria fridge nicely to keep your food fresh for the day. We may be familiar with such behaviour from cases where pet-owners insist that everyone participate in the fiction that their pet is just as entitled to moral consideration as humans, and thus is owed a literal place at the table as well as the same three-course meal as the human dinner guests.

Although it may seem intuitively plausible that something highly problematic is going on here, it may not yet be clear what makes it degrading to treat non-members of the moral community as if they were members, according to the relational approach to dignity. As mentioned in Section 4, the standard way of accounting for offence to dignity relies on *other* members treating one in disrespectful ways; but in the case of chatbot interaction,

no other members of the moral community are directly involved. So why would our dignity be at issue in such cases?

Recall that second-personal address is inherently reciprocal in nature. Both in levying and in acknowledging a moral claim, one does not just assert one's own standing, but that of the other party as well, because a normative claim of this kind can only be understood, and hence acted upon, by a being that is itself a moral agent. Reciprocity is crucial to dignity. Since recognizing another person's second-personal authority is to acknowledge that person has normative authority *over* one, a unilateral acknowledgement of such standing would amount to an act of submission. It would signify that one takes oneself to have a lesser standing than the entity that wields such authority over one, but grants one no such authority over it in return. Without reciprocity, ascribing second-personal authority to an entity is a form of self-debasement.

The case of Incitatus illustrates this clearly. Imagine that the senators were to try – as best as they can – to acknowledge that Incitatus has second-personal standing. They would, for instance, treat him with all the attitudes and accoutrements we typically reserve for a fellow moral agent, thus implying that Incitatus can levy moral claims on them. For this to be compatible with their own dignity, Incitatus would have to respond in kind. But it is obvious that, due to his equine nature, Incitatus cannot do so. As a result, the senators' acknowledgement remains one-sided: they have, in their behaviour, signalled that Incitatus has second-personal authority over them, but they have no such authority over Incitatus, as Incitatus is not able to understand second-personal authority.

We submit that the case is similar for AI systems. Though chatbots are undoubtedly much better at mimicking human interaction, and thus may give the *appearance* of reciprocating in kind, they lack the capacities to respond to moral reasons and form the relevant attitudes necessary for moral agency. As a result, they cannot reciprocate second-personal respect. Thereby, treating them as if they were moral agents affronts our dignity. No one will be fooled by Incitatus, as he does not, and cannot, even seem to show second-personal respect. Chatbots, however, have become very good at producing human-like outputs, typical marks of moral recognition and respect included. This superficial appearance of recognition and respect invites and encourages behaviour that goes against our duties of self-respect, placing users in an asymmetrical relation where they are to (implicitly) recognize or signal moral authority where there is none, while not receiving any such reciprocal recognition in return.

Let us now return to our chatbot use cases and discuss in more detail how they can instantiate such threats to users' dignity.

a) Chatbots as Information Retrievers and Task Automators

Using chatbots as information retrievers and for automating some tasks, such as grammar correction and translation, seems rather innocuous from the point of view of relational theories of dignity. When users issue pure commands, as they would do with an ATM or automated food dispenser, they do not presuppose moral or rational attitudes by the technology they interact with. Consequently, there seems to be no risk to their dignity at play.

However, such tasks vary in complexity, and it is easy to slip into requests that presuppose more from chatbots than a mere ability to retrieve the information desired, or to automate a task that would otherwise be tedious. Consider two sets of contrasting examples:

Innocent Retrieval: Tell me the weather forecast for Oslo for the next two hours.

Innocent Automation: I need a workout schedule that spreads my activities evenly over the week. I need three sets of exercises: conditioning, strength, and flexibility. Provide some suggestions to fit these activities into the attached weekly schedule.

Slippery Retrieval: I'm confused about the political situation in the Middle East. What are the most reliable sources covering the recent developments and its background history?

Slippery Automation: I want to increase my productivity at work, but I do not want to neglect my family and friends, and need to sustain my mental health. Revise the attached weekly schedule to increase my work productivity according to these demands.

At first glance, the Slippery and the Innocent cases appear similar: revising a weekly schedule and providing information about the world. In the Slippery cases, however, executing these tasks requires complex moral and rational attitudes: what decrease in time with family and friends constitutes 'neglect'; which of the many sources on the political situation in the Middle East are reliable? By issuing such requests to chatbots, one effectively presupposes that they have the requisite abilities and attitudes to give the matter proper consideration, but they certainly lack those. They are trained on large datasets reflecting a wide variety of attitudes on such topics as the Middle East and work-life balance, and execute a complex probabilistic calculation based on those datasets to produce an output. These calculations do not incorporate a moral attitude about what 'neglect' consists of, nor a normative epistemological attitude as to what counts as reliable reporting.²⁸

We can see the importance of such attitudes clearly when we consider human interactions. Which friends or acquaintances would you turn to regarding the slippery cases? For information about the Middle East, you would ask someone you trust to have a critical and informed attitude towards news reporting, history, and politics. For information on how to maintain a healthy work-life balance, you would ask someone you know to be a responsible employee and dedicated friend, partner, and parent. Such sources are much more trustworthy than an acquaintance who has encyclopaedic knowledge of all that has been written on these issues but lacks a normative compass altogether. We might ask more carefully curated questions to this latter acquaintance, along the lines of the innocent cases above. They can report on what is out there, but we do not rely on them to do the requisite selection in providing the response we seek.

Some tasks involving merely information retrieval and automation require moral and rational attitudes that are absent in chatbots. When relying on chatbots to execute those tasks, we are implicitly ascribing such attitudes to objects that clearly lack them. The relational account of dignity we espouse here indicates that such requests to chatbots offend the dignity of those making them.

b) Chatbots as Helpdesk Service

As with the information retrieval case, there are innocuous uses of chatbots in helpdesk services. You want to rebook a flight, and the chatbot provides a link to the rebooking page. You consider switching phone plans, and the chatbot provides a detailed overview of the available options. Your train is delayed, and the chatbot informs you that you will

still make your connection. In such cases, helpdesk chatbots function as no more than information retrieval services.

But helpdesk chatbots are used for other purposes too. One may contact a company to lodge a complaint about the services delivered, and be made to address it to a chatbot. Here, dignity is more likely at issue, as complaining is inherently second-personal. A complaint does not just seek to remedy whatever the complaint is about, but also seeks second-personal recognition. What it seeks to evoke, at least in part, is the confirmation that the other party has the right attitude. A complaint seeks remedy, but also acknowledgement that a wrong has occurred – and such acknowledgement must be sincere to have value.

To see this, we can imagine that the chatbot provides the same remedies that would be offered by a compassionate human, say, a rebooking or a partial refund. Moreover, we can imagine that the remedies are accompanied by the strings of words we would expect from a commiserating human: ‘I’m very sorry that this happened’, or ‘we hope this will help you with your problem’. What is certainly missing in the chatbot case is an actual moral attitude towards the situation you have found yourself in, or a rational attitude as to whether it was indeed worthy of complaint.

Often enough, these attitudes are not in place as we would hope in human interactions either – the stereotypical helpdesk worker reiterates company policy in a robot-like fashion, something that can be disrespectful in its own right – but human interlocutors at least *are in a position* to entertain the moral and rational attitudes mentioned above, rather than merely taking you or your complaint as nothing more than a problem to solve. And indeed, experience teaches that stereotypes know plenty of exceptions; helpdesk workers are often helpful, friendly, and ready to acknowledge that the company in question ruined your day. Chatbots do not allow for any such exceptions: even when their outputs have the marks of friendliness and acknowledgement, they are not expressions of such attitudes. Chatbots can *manage* complaints in the thinnest sense of the word, but they cannot respect the user as required by the practices of accepting and remedying complaints by members of the moral community.

c) Chatbots as Advisors

Some ways of using chatbots for advice appear innocuous and indeed equivalent to the unproblematic use of these systems as mere information retrievers. Finding oneself in a pickle, one may ask a chatbot for information on how similar circumstances have been tackled by others, or tips that they may have given. The chatbot, in such a case, would work as a sort of life-FAQ, provided that its outputs are neutral and impersonal as we would expect from such a document.

However, this is often not the case, with chatbot outputs seemingly providing pieces of advice directly, and promoting certain courses of action in their own name, as it were; rather than listing them as third-personal information recompiled, sometimes only with the slightest changes, in the output.²⁹ For instance, if a user should ask which aspects of managing family conflicts are most important, they may receive an output such as: ‘I emphasise effective communication as the most important aspect of resolving family conflicts because it serves as the foundation for understanding and collaboration’.³⁰

Outputs of this kind are problematic in several respects. First, they mislead the user about the nature of the chatbot, implying that it has attitudes and capacities that it does not in fact possess, such as the capacity to judge and to emphasize. As such, the outputs

suggest that the chatbot is an agent capable of attitudes and capacities that would make it, to a certain degree at least, an interlocutor on a par with the user. Second, and connectively, the content of the outputs and the recommendation of potential continuations for the interaction suggest that it is sensible to ask the chatbot for life advice. The sort of interaction that is encouraged is thereby much more akin to that between persons than to that between a person and an FAQ document, an ATM, or a piece of text-editing software.

Finally, and most importantly, chatbot outputs such as these encourage persons to seek and accept advice from a non-member of the moral community. While one may follow instructions shown on the screen of an ATM, no one will be inclined to regard the ATM as an agent that is providing one with advice or commands. The machine-external source of the instructions is evident to users, as much as when they are written on the user manual of a domestic appliance.

Advice-giving and -receiving is a practice that requires the kinds of mutual recognition and respect that are only available to members of the moral community. To give and receive advice requires a recognition of the other as someone with experiences, rational and emotional attitudes, and moral standing, and thereby entitled to special consideration of their needs and demands, as well as obligations. In addition, the advice-giver typically bears some responsibility for the quality of the advice, becoming, for instance, liable to reactive attitudes like praise, gratitude, blame, or resentment, depending on whether the advice proves sound. Current chatbots do not have what it takes to be participants in such practices. To encourage persons to engage in advice-receiving practices with chatbots is thus to invite a kind of self-debasement, a violation of self-respect. It involves taking seriously guidance and advice from something that does not and cannot recognize one as a person, and that cannot engage sincerely in such practice. The veneer of linguistic competence that marks chatbots disguises the fact that asking them for advice is akin to asking an ATM for life tips.

If the user manages to resist the temptation to anthropomorphize the chatbot, treating it as nothing more than a fancy information retrieval engine, the ethical risk might be mitigated. However, as pointed out, humans are very sensitive to anthropomorphic cues. Being partly designed to display such cues abundantly, it is unclear to what extent one can keep at bay such a temptation, both consciously and subconsciously. Be this as it may, such design choice incentivizes users to behave in the aforementioned self-degrading ways. Such incentives would be ethically problematic even in an idealized world where all chatbot users were able, through their own stalwart effort, to withstand them.

d) Chatbots as Companions

Similar considerations apply to the use of chatbots as companions, that is, as systems designed to play the role of friends, family members, or romantic partners. In this context, the dignity-offending interactions involve a host of practices that can only be truly engaged in by members of the moral community. In this way, companion-chatbots are the most extreme extant case of the kind of failure of self-respect that we have been belabouring in this article.

Indeed, companion-chatbots are supposed to display and elicit emotional attachment; simulate the dynamics of interpersonal recognition, care, and trust; as well as to function as emotional supporters, confidants, and advice-givers. Such systems invite and

encourage their users to violate the demands of self-respect in much the same way as described in the previous cases, only more egregiously so.

Not only is the range of domains over which human behaviour is superficially mimicked much larger, but those domains are ones in which reciprocity and mutual recognition as moral, rational beings are arguably even more central. In the best-case scenario, a user may treat a chatbot's piece of advice as little more than an (ethically problematic) misleading way of providing third-personal information. It is hard to see how this could be the case with companion-chatbots.

Companion-chatbots are typically designed and marketed as AI companions. The company Replika, for instance, presents its chatbot as 'an AI companion who is eager to learn and would love to see the world through your eyes ... Replika is always ready to chat when you need an empathetic friend'.³¹ Its website includes testimonials in which users claim to have developed feelings for the chatbot, ascribing a variety of human capacities and features to it, such as being cheerful, supportive, available, and comforting.³² In brief, the whole point of AI companions is to treat them as if they were members of the moral community, so as to be given in return the simulacrum of reciprocated sympathy, care, support, friendship, respect, and love.

It is likely that, for many users, engaging with such chatbots is a conscious exercise in make-believe. This does not, though, diminish the fact that such systems are not presented as platforms for make-believe exercises, but rather as systems that can authentically provide things such as emotional support and friendship. This is in itself ethically problematic, regardless of whether users take such claims seriously. Moreover, the abundance of anthropomorphic cues indicating agency, emotion, mentality, and so on, are extremely difficult to resist. Even savvy users may therefore partially, or occasionally, fall prey to the illusion of interacting with a moral agent.³³ This is made (even) more pressing by recent technological developments, which enable chatbot outputs to be conveyed by human-like synthetic voices, and the chatbots to be represented through avatars that display context-appropriate facial and bodily expressions and cues. The development has been, in other words, towards *more* human-likeness, superficial as it may be, rather than less.

There is evidence that AI companions can be useful tools for tackling loneliness, social isolation, and other factors that negatively influence wellbeing.³⁴ Our concern in this article is orthogonal to such considerations. Even if the use of AI companions should prove beneficial for some people in certain circumstances, it would still involve an offence to the standing of moral beings – and it should be noted that the willingness to put one's interests above the demands of one's dignity is a paradigmatic instance of a violation of self-respect.

6. Concluding Remarks

In this article, we have argued that advanced chatbots pose an ethical risk that has not been adequately appreciated in current literature on the ethics of AI-based conversational systems. In particular, we proffered that such chatbots offend users' dignity by encouraging anthropomorphic treatment, thereby inviting users to treat and interact with them in ways appropriate only to something they are not, namely members of the moral community.

This leads to the impossibility of the mutual recognition of moral standing that is constitutive of dignity, and with it a failure of self-respect, a form of moral debasement.

We discussed four common chatbot use cases to illustrate the partly different ways in which they represent threats to the dignity of users, going from the relatively unproblematic – chatbots as pure information retrievers – to the deeply problematic, that is, chatbots as companions. The human-likeness of current chatbot outputs may appear to have advantages, as it can make interactions seem more natural and user-friendly. This, however, comes at a high ethical cost. Chatbots should not thus be designed to emulate features of human agents, and should not be used as replacements in practices that involve the sort of mutual recognition that comes with being members of the moral community.

The list of ethical risks associated with chatbots is already very long.³⁵ In this article, we have made it a little bit longer.

Jan-Willem van der Rijt, Dimitri Coelho Mollo and Bram Vaassen, Umeå University, Umeå, Sweden. jan-willem.vanderrijt@umu.se

Acknowledgements

We wish to thank Björn Lundgren and Nancy Nuñez, whose invitation to their workshop ‘AI and Human Dignity’ (Erlangen, 2024) prompted us to write this article. We are grateful to the participants of this workshop, the referees of the *Journal of Applied Philosophy*, as well as audiences in Linköping and Gothenburg for their helpful comments.

NOTES

- 1 See e.g. Weidinger et al., “Ethical and Social Risks”; Gabriel et al., “Ethics.”
- 2 There is, however, much less agreement on what grounds such special standing, a debate we cannot go into here (but see e.g. Debes, *Dignity*; Düwell et al., *Cambridge Handbook*; McCrudden, *Understanding*; Rosen, *Dignity*; Schaber, *Menschenwürde*; von der Pfordten, *Menschenwürde*, for overviews).
- 3 Cf. Darvall, *Second-Person Standpoint*; Zylberman, “Two Second-Personal Conceptions.”
- 4 On one reading, such a response is mere cajolement, which would be problematic enough, but on another reading the chatbot appears to be insisting on a justification, which, as we argue below, amounts to arrogating to itself a status it simply does not have.
- 5 Microsoft CoPilot, accessed on September 30, 2024.
- 6 Microsoft CoPilot, accessed on September 30, 2024.
- 7 Such applications make up a substantial part of overall chatbot use. Recent reports estimate some 1.4 billion users worldwide make use of AI companions (Ferguson, “AI Chatbots”).
- 8 See Skjuve et al., “My Chatbot Companion.”
- 9 See Weizenbaum, *Computer Power*.
- 10 This human-likeness is only partial, though; see McCoy et al., “Embers,” for examples of failures.
- 11 Cf. Shevlin and Halina, “Apply Rich Psychological Terms”; Zimmerman et al., “Human/AI Relationships”; Shanahan, “Talking.”
- 12 Some argue that certain AI systems may soon acquire, or even already have, moral status (e.g. Altehenger et al., “How AI Systems”). While we cannot settle this issue here, we note that these authors typically operate with a thinner notion of moral status than we do, and that the chatbots under consideration here would not clear even their lowered bar for moral status.
- 13 Cf. Weidinger et al., “Ethical and Social Risks”; Dahlgren Lindström et al., “Helpful”; Gabriel et al., “Ethics.”
- 14 See Kim and Sundar, “Anthropomorphism”; Salles et al., “Anthropomorphism”; van der Goot et al., “Understanding”; Gabriel et al., “Ethics.”

- 15 Kant, "Metaphysics."
- 16 See e.g. Hill, "Servility and Self-Respect"; Hill, "Self-Respect Reconsidered"; Dillon, "How to Lose."
- 17 See Darwall, *Second-Person Standpoint*.
- 18 Ibid., 5ff.
- 19 For the close link between dignity and the practice of demanding or claiming, see also Feinberg, "Nature."
- 20 Darwall, *Second-Person Standpoint*, 8, 20ff.
- 21 See ibid., 3, 15–7; cf. Strawson, "Freedom."
- 22 Margalit, *Decent Society*, chap. 7.
- 23 We rely here on a broadly Kantian–Darwallian notion of dignity because we believe it to be the most convincing account of dignity available at present (for a partial defence, see Van der Rijt, "Inherent Dignity"; Van der Rijt, "Transmutation"); its superior ability to explain violations of dignity being a particularly weighty consideration in its favour. However, we expect most alternative accounts that also embrace the special nature and/or importance of human beings or of moral agency to be broadly sympathetic to the case we make in this article. Those who downplay or deny the special importance of moral personhood, on the other hand, and wish, for instance, to extend dignity throughout the animal kingdom (e.g. Nussbaum, *Creating Capabilities*; or Gilibert, "Inclusive Dignity"), might find our general line of argument less congenial.
- 24 Cf. e.g. Schaber, *Instrumentalisierung*.
- 25 Cf. Hill, "Servility and Self-Respect."
- 26 We limit our claims to current chatbots as we wish to remain neutral on the question of whether future AI systems could come to possess the rich degree of moral agency required to count as full members of the moral community (i.e. to effectively achieve full moral personhood). At any rate, our arguments are meant to apply to currently existing chatbots as well as any future chatbots relevantly similar to current ones, i.e. that are endowed with analogous capabilities and that display comparable behaviour.
- 27 See Lopatovska and Williams, "Personification"; Skjuve *et al.*, "My Chatbot Companion."
- 28 The strongest sense in which chatbots can be said to incorporate any such attitudes is that they produce strings of text reflecting the dominant attitudes in the dataset (*modulo* hand-coded constraints imposed by the developers). Requesting information from such an amorphous blob comes with its own ethical and epistemological challenges that we will not address here, but see Dahlgren Lindström *et al.*, "Helpful," for a detailed analysis.
- 29 Microsoft CoPilot sometimes, but not always, appends links to the end of sentences. Other than that (and general service disclaimers), there is often no further signalling, unless actively sought by the user, that the advice given is not the chatbot's own.
- 30 Microsoft CoPilot, accessed October 1, 2024.
- 31 Replika.com, accessed October 7, 2024. The website also publicizes the fact that the user can choose whether the chatbot is supposed to 'be' a friend, a partner, or a mentor. Note also the misuse of the personal pronoun 'who', rather than the grammatically correct and morally appropriate 'that'.
- 32 Skjuve *et al.*, "My Chatbot Companion."
- 33 Skjuve *et al.* (ibid.) describe in detail how human–AI companion 'relationships' tend to develop from an initial scepticism towards deep forms of emotional attachment.
- 34 Ibid.
- 35 See Gabriel *et al.*, "Ethics."

References

- Altehenger, H., L. Menges, and P. Schulte. "How AI Systems Can be Blameworthy." *Philosophia* 52 (2024): 1083–1106.
- Dahlgren Lindström, A., L. Methnani, L. Krause, P. Ericson, Í. M. de Rituerto de Troya, D. Coelho Mollo, and R. Dobbe. "Helpful, Harmless, Honest? Sociotechnical Limits of AI Alignment and Safety through Reinforcement Learning from Human Feedback." *Ethics and Information Technology* 27, no. 2 (2025): 28.
- Darwall, S. *The Second-Person Standpoint: Morality, Respect and Accountability*. Cambridge: Harvard University Press, 2006.
- Debes, R., ed. *Dignity: A History*. Oxford: Oxford University Press, 2017.
- Dillon, R. "How to Lose your Self-Respect." *American Philosophical Quarterly* 29 (1992): 125–139.

- Düwell, M., J. Braarvig, R. Brownsword, and D. Mieth, eds. *The Cambridge Handbook of Human Dignity*, eds. ed. Cambridge: Cambridge University Press, 2014.
- Feinberg, J. "The Nature and Value of Rights." *Journal of Value Inquiry* 4 (1970): 243–257.
- Ferguson, M. "AI Chatbots: the Future of Companionship or Cause for Concern?" 2024. OpenTools. <https://opentools.ai/news/ai-chatbots-the-future-of-companionship-or-cause-for-concern#section3>.
- Gabriel, I., A. Manzini, G. Keeling, L. A. Hendricks, V. Rieser, H. Iqbal, et al. "The Ethics of Advanced AI Assistants." 2024. <https://arxiv.org/abs/2404.16244>.
- Gilbert, P. "Inclusive Dignity." *Politics, Philosophy & Economics* 23 (2024): 22–46.
- Hill, T. "Servility and Self-Respect." In *Dignity, Character, and Self-Respect*, edited by R. Dillon, 76–92. Abingdon: Routledge, 1996a.
- Hill, T. "Self-Respect Reconsidered." In *Dignity, Character, and Self-Respect*, edited by R. Dillon, 117–124. Abingdon: Routledge, 1996b.
- Kant, I. "The Metaphysics of Morals." In *The Cambridge Edition of the Works of Immanuel Kant: Practical Philosophy*, 353–604. Cambridge: Cambridge University Press, 1797/1996.
- Kim, Y., and S. S. Sundar. "Anthropomorphism of Computers: Is it Mindful or Mindless?" *Computers in Human Behavior* 28, no. 1 (2012): 241–250.
- Lopatovska, I., and H. Williams. "Personification of the Amazon Alexa: BFF or a Mindless Companion." In *Proceedings of the 2018 Conference on Human Information Interaction & Retrieval – CHIIR '18*, 265–68. New York: ACM Press, 2018.
- Margalit, A. *The Decent Society*. Cambridge: Harvard University Press, 1996.
- McCoy, R. T., S. Yao, D. Friedman, M. D. Hardy, and T. L. Griffiths. "Embers of Autoregression Show How Large Language Models are Shaped by the Problem They are Trained to Solve." *Proceedings of the National Academy of Sciences* 121, no. 41 (2024).
- McCrudden, C., ed. *Understanding Human Dignity*. Oxford: Oxford University Press, 2014.
- Nussbaum, M. *Creating Capabilities*. Cambridge: Harvard University Press, 2011.
- Rosen, M. *Dignity: Its History and Meaning*. Cambridge: Harvard University Press, 2012.
- Salles, A., K. Evers, and M. Farisco. "Anthropomorphism in AI." *AJOB Neuroscience* 11, no. 2 (2020): 88–95.
- Schaber, P. *Instrumentalisierung und Würde*. Paderborn: Mentis, 2010.
- Schaber, P. *Menschenwürde*. Stuttgart: Reclam, 2012.
- Shanahan, M. "Talking about Large Language Models." *Communications of the ACM* 67, no. 2 (2024): 68–79.
- Shevlin, H., and M. Halina. "Apply Rich Psychological Terms in AI with Care." *Nature Machine Intelligence* 1 (2019): 165–67.
- Skjuve, M., A. Følstad, K. I. Fostervold, and P. B. Brandtzaeg. "My Chatbot Companion – a Study of Human–Chatbot Relationships." *International Journal of Human–Computer Studies* 149 (2021): 102601.
- Strawson, P. "Freedom and Resentment." In *Freedom and Resentment and Other Essays*, 1–28. Abingdon: Routledge, 1974/2008.
- van der Goot, M. J., N. Koubayová, and E. A. van Reijmersdal. "Understanding Users' Responses to Disclosed vs. Undisclosed Customer Service Chatbots: A Mixed Methods Study." *AI and Society* 39, no. 6 (2024): 2947–60.
- van der Rijt, J. "Inherent Dignity, Contingent Dignity and Human Rights: Solving the Puzzle of the Protection of Dignity." *Erkenntnis* 182 (2017): 1321–38.
- van der Rijt, J. "The Transmutation of Dignity: Kant, Neo-Roman Republicanism, and the Commonwealth of Ends." In *Human Dignity and the Kingdom of Ends*, edited by J. van der Rijt and A. Cureton, 148–167. Abingdon: Routledge, 2022.
- von der Pfordten, D. *Menschenwürde*. Munich: Verlag C.H. Beck, 2016.
- Weidinger, L., J. Mellor, M. Rauh, C. Griffin, J. Uesato, P.-S. Huang, M. Cheng, et al. "Ethical and Social Risks of Harm from Language Models." 2021. <https://arxiv.org/abs/2112.04359>.

- Weizenbaum, J. *Computer Power and Human Reason*. San Francisco: Freeman, 1976.
- Zimmerman, A., J. Janhonen, and E. Beer. "Human/AI Relationships: Challenges, Downsides, and Impacts on Human/Human Relationships." *AI and Ethics* 4, no. 4 (2023): 1555–67.
- Zylberman, A. "Two Second-Personal Conceptions of the Dignity of Persons." *European Journal of Philosophy* 25, no. 4 (2017): 921–943.