

When “AI Understands” Becomes Safe to Say: Anthropomorphic Language, Metaphor Stabilization, and the Cultural Acceptance of Non-Subjective Dialogue Systems

Karel Hrubec

Independent researcher

ORCID: 0009-0005-9569-5662

DOI: 10.5281/zenodo.21056975

Working manuscript v0.3 — June 2026

Abstract

This conceptual paper argues that the central risk of anthropomorphic language in contemporary artificial intelligence is not anthropomorphism itself, but the absence of a culturally stabilized metaphorical frame. Expressions such as “the engine does not want to start” are normally harmless because their metaphorical status is socially settled: users do not infer intention, frustration, or subjective experience from the phrase. In contrast, expressions such as “AI understands,” “AI thinks,” or “from my perspective” remain unstable because dialogue, contextual responsiveness, and first-person linguistic forms have historically been associated with subjectivity. The paper proposes that contemporary AI systems occupy a missing cultural category: *non-subjective dialogue systems* capable of operating with language, meaning, probability, and context without thereby becoming subjects of experience. Building on the ELIZA effect, the Computers Are Social Actors / Media Equation tradition, and recent work on LLM anthropomorphism, the paper reframes AI acceptance as a process of *metaphor stabilization* under repeated interaction. In response to expected objections, the paper explicitly distinguishes descriptive claims about user perception, normative claims about responsible communication, and metaphysical claims about consciousness that it does not attempt to settle. It also proposes an operationalization of metaphor stabilization and a minimal pilot protocol for empirical testing. The paper concludes that responsible AI communication should neither erase anthropomorphic shorthand nor intensify it uncritically, but should help users stabilize the distinction between dialogic form and subjective ownership of dialogue.

Keywords: AI anthropomorphism; metaphor stabilization; non-subjective dialogue systems; ELIZA effect; media equation; relational objectivity; AI acceptance; human-AI interaction; artificial intelligence; subjectivity

1 Introduction: The Unsafe Metaphor

A person may say:

The engine does not want to start.

The sentence is technically inaccurate. Engines do not want, resist, intend, or become frustrated. Yet the expression is normally harmless. It functions as a culturally stabilized metaphor: the speaker and listener understand that “does not want” means something like “does not perform the expected function under the current operating conditions.”

Now consider a similar expression:

The AI does not want to answer.

Or:

The AI understands the question.

Or:

From my perspective, the answer is . . .

These expressions are not merely equivalent to the engine metaphor. They occur in a different cognitive and cultural space. The engine metaphor is safely read as figurative because the engine is not linguistically or socially close to a human subject. Dialogue-based AI systems, especially large language models (LLMs), enter through precisely the channel historically associated with subjectivity: language, response, explanation, memory-like context, apparent self-reference, and dialogic adaptation.

The problem is not that anthropomorphic language exists. Human language is full of anthropomorphic shorthand. The problem is that anthropomorphic language becomes risky when its metaphorical status is not yet socially stabilized.

This paper argues that contemporary AI exposes a missing category in ordinary human understanding. AI is not a human subject, but it is also not a simple tool in the traditional sense. It can participate in dialogue without being a subject of dialogue. It can generate the linguistic form of perspective without owning a perspective. It can operate with meaning-like structures without thereby becoming a living bearer of meaning.

The central thesis is therefore narrow:

Anthropomorphic language becomes relatively safe when it is culturally understood as metaphor. In the case of AI, this stabilization has not yet fully occurred.

This paper does not claim that AI anthropomorphism is new. The ELIZA effect, the Computers Are Social Actors (CASA) paradigm, and the Media Equation tradition already show that humans readily apply social and mentalistic interpretations to computers ([Weizenbaum, 1966](#); [Nass and Moon, 2000](#); [Reeves and Nass, 1996](#)). The contribution of this paper is different. It proposes that the cultural acceptance of contemporary AI requires stabilizing a new everyday category: the *non-subjective dialogue system*.

2 Scope and Claim Boundaries

Because anthropomorphism, consciousness, and AI acceptance are easily conflated, the argument must begin with explicit boundaries. This paper makes three kinds of claims and deliberately avoids a fourth.

2.1 Descriptive Claim

The descriptive claim is that users may interpret AI’s dialogic behavior through inherited categories such as person, mind, assistant, tool, or program. These categories are partially useful but unstable when applied to systems that generate subject-like language without demonstrated subjectivity. This instability is not a rare user failure; it is a predictable response to linguistic and social cues.

2.2 Conceptual Claim

The conceptual claim is that a missing category is needed: *non-subjective dialogue system*. This category names a system that can participate in meaningful linguistic interaction without thereby being a subject of experience, intention, or existential concern.

2.3 Normative Claim

The normative claim is that responsible AI communication should help users preserve the distinction between dialogic form and subjective ownership of dialogue. This does not require banning first-person language or conversational shorthand. It requires making the metaphor stabilizable and clarifiable, especially in high-stakes or emotionally vulnerable contexts.

2.4 Metaphysical Non-Claim

The paper does *not* claim to solve the metaphysics of machine consciousness. It does not prove that no artificial system could ever be conscious. It only states that present dialogue systems should not be treated as subjects merely because they can generate the linguistic form of subjectivity. In other words, the paper is operational and cultural rather than metaphysically final.

3 Background: Anthropomorphism Is Not New

The tendency to attribute human-like understanding to computational systems is not a recent discovery. Joseph Weizenbaum’s ELIZA demonstrated that even relatively simple pattern-matching systems could evoke in users the impression of understanding (Weizenbaum, 1966). ELIZA did not understand in any human sense. Yet its conversational structure was sufficient to invite projection.

This phenomenon later became known as the ELIZA effect: the tendency to project comprehension, empathy, intentionality, or personality onto computational outputs. The effect is not merely a technical misunderstanding. It reveals something about human cognition: when

a system produces language in a socially recognizable form, users often complete the missing subject behind the utterance.

The CASA paradigm and the broader Media Equation tradition extended this insight. Humans often respond to computers and media using social rules, even when they explicitly know that the system is not human (Nass and Moon, 2000; Reeves and Nass, 1996). Politeness, reciprocity, personality attribution, and social evaluation can be triggered by relatively minimal cues.

Contemporary LLMs intensify the issue. They do not merely produce simple scripted responses. They can sustain long conversations, adapt tone, explain themselves, respond to emotional content, and use first-person language. Recent studies on LLM anthropomorphism show that modality, grammatical person, perceived warmth, competence, and empathy can influence how users attribute human-like qualities, trust, relational closeness, usefulness, and perceived accuracy to AI systems (Cohn et al., 2024; Kadambi et al., 2026). Related work proposes measurement approaches for anthropomorphism in language (Cheng et al., 2024), investigates anthropomorphized system descriptions (Inie et al., 2024), evaluates anthropomorphic LLM behaviors in multi-turn interactions (Ibrahim et al., 2025), and treats anthropomorphism as a design variable rather than as a single unavoidable default (Xiao et al., 2025; So et al., 2026; Ferrario et al., 2026).

This literature shows that anthropomorphism is a known and active research area. Therefore, the present paper must avoid an overclaim. Its claim is not:

Humans anthropomorphize AI.

That is already known. The narrower claim is:

AI anthropomorphism is unusually unstable because society has not yet stabilized the metaphorical status of AI’s subject-like language.

4 Terminology

Table 1 defines the core terms used in the paper. These definitions are intentionally operational. Their purpose is not to settle all philosophical disputes, but to make the proposed framework testable and criticizable.

Table 1: Operational terminology.

Term	Operational definition
Anthropomorphic language	Language that attributes or implies human-like mental, affective, social, or agential properties to a non-human system. Examples include “AI thinks,” “AI wants,” “AI understands,” and first-person self-reference.

Term	Operational definition
Dialogic form	The observable linguistic form of dialogue: turn-taking, response to meaning, contextual continuation, explanation, apparent self-reference, and adaptive tone.
Subjective ownership of dialogue	The condition in which dialogue is produced by a subject who has first-person experience, existential stakes, affective involvement, and ownership of meaning. This paper does not claim that current AI systems possess this.
Non-subjective dialogue system	A system capable of participating in linguistically meaningful interaction without thereby being a subject of experience, intention, or existential concern.
Metaphor stabilization	The process by which a community learns to use a figurative expression without routinely mistaking it for an ontological claim.
Unstable anthropomorphic cue	A cue that may be read either as metaphorical shorthand or as evidence of actual subjectivity, because its status is not yet socially settled.
Relational acceptance	The gradual stabilization of a technology’s meaning through repeated interaction, correction, failure, reliance, education, and institutional framing.

The terminology is intentionally somewhat restrictive. Existing labels such as *tool*, *software*, and *chatbot* often under-describe the dialogic and semantic capacity of contemporary systems. Conversely, labels such as *agent*, *assistant*, *companion*, or *collaborator* may over-invite agency, social role, or subject-like interpretation. The term *non-subjective dialogue system* is therefore not proposed as an elegant public-facing label, but as an analytically protective category: it preserves the dialogue function while blocking premature attribution of subjective ownership.

Similarly, *subjective ownership of dialogue* is used instead of the simpler term *consciousness* because the paper does not attempt to solve consciousness metaphysics. The relevant distinction is narrower and communicative: whether the system merely generates the form of dialogue, or whether dialogue belongs to a subject for whom meaning is lived, affectively weighted, and existentially situated.

5 The Missing Category

Ordinary language offers several familiar categories for artificial systems:

- tool,
- machine,
- program,

- robot,
- assistant,
- agent,
- mind,
- person.

None of these is fully adequate for contemporary dialogue-based AI.

Calling AI a “tool” is partly correct but incomplete. A hammer does not respond to the semantic content of a request. A calculator operates with numbers according to determinate rules, but it does not sustain open-ended dialogue. Classical software executes predefined procedures, but it does not normally generate context-sensitive explanations in natural language.

Calling AI a “person” is also misleading. Present AI systems do not have demonstrated subjective experience, embodied vulnerability, existential concern, biological self-maintenance, or first-person ownership of meaning. They can simulate first-person language, but simulation of first-person language is not equivalent to being a first-person subject.

This paper therefore proposes the category:

Non-subjective dialogue system

A non-subjective dialogue system is a system capable of participating in linguistically meaningful interaction without thereby being a subject of experience, intention, or existential concern.

The category is deliberately minimal. It does not solve the metaphysics of consciousness. It does not prove that no artificial system could ever become conscious. It only provides an operational distinction for current AI communication:

Dialogic form is not identical with subjective ownership of dialogue.

This distinction matters because contemporary AI often appears in the social position formerly occupied almost exclusively by persons. It answers questions, gives advice, uses “I,” apologizes, explains its limitations, and appears to understand. The user is therefore exposed to a form that historically implied a subject, while the system itself does not provide evidence of subjectivity.

The result is category instability.

6 Metaphor Stabilization

The key concept proposed here is *metaphor stabilization*.

Metaphor stabilization is the process by which a community learns to use a figurative expression without routinely mistaking it for an ontological claim.

The phrase “the engine does not want to start” is stabilized. Most speakers understand that “want” is not being used literally. The engine has no preference. The sentence functions as shorthand for a functional failure under given conditions.

The phrase “AI understands” is not yet equivalently stabilized. Some users treat it as shorthand for a system generating contextually appropriate outputs. Others treat it as evidence of comprehension. Others remain uncertain. The same phrase oscillates between metaphor, functional description, and ontological implication.

This instability is not caused only by user ignorance. It is also caused by the design of the interaction. AI systems often use the first person because ordinary conversation becomes more fluid that way. A system may say:

I think the best answer is ...

A more technically precise version would be:

Based on the available input, context, and learned patterns, the generated response favors the following interpretation ...

The second formulation is ontologically cleaner, but communicatively burdensome. The first is efficient, but anthropomorphically suggestive.

Therefore, the issue is not whether anthropomorphic shorthand should be absolutely forbidden. Such a prohibition would make many interactions awkward and inefficient. The issue is whether the system and the surrounding culture provide enough stabilization for users to understand what the shorthand means.

A stabilized AI metaphor would allow users to say “AI understands” in the same limited way that they say “the engine does not want to start.” It would mean:

The system generated a contextually appropriate response that tracks the relevant semantic relations.

It would not mean:

The system possesses subjective understanding.

Until this distinction becomes culturally stable, anthropomorphic AI language remains more dangerous than ordinary machine metaphors.

7 Operationalizing Metaphor Stabilization

One criticism of a purely conceptual account is that “metaphor stabilization” may remain suggestive but vague. This section proposes measurable indicators.

A community, user group, or experimental participant can be said to exhibit stronger metaphor stabilization when the following conditions hold:

1. **Low subjective attribution:** users do not infer consciousness, feeling, personal intention, or existential concern from conversational shorthand.

2. **High functional interpretation:** users correctly translate statements such as “AI understands” into operational descriptions such as “the system generated a contextually appropriate response.”
3. **Failure-resistant interpretation:** when the system fails, users do not explain the failure primarily through human-like motives such as unwillingness, offense, laziness, resentment, or deception unless deception is specifically evidenced at the organizational or design level.
4. **Clarification recall:** users remember and can restate that first-person language is a conversational convention rather than a claim of subjectivity.
5. **Stable trust calibration:** users can maintain appropriate trust in useful outputs without increasing belief in subjective experience.

These indicators make the concept testable. They also distinguish metaphor stabilization from mere distrust. A user who rejects all AI outputs because the system is “just a machine” has not necessarily stabilized the metaphor. Stabilization requires the middle category: the system can be operationally useful and dialogically competent without being a subject.

8 AI Acceptance as Relational Stabilization

Technology acceptance is often discussed in terms of usefulness, trust, usability, or perceived risk. These dimensions are important, but they do not fully capture the cultural difficulty of AI. Contemporary AI is not only a tool to be adopted. It is also a new kind of interaction partner to be conceptually placed.

The public does not yet possess a stable everyday category for a system that behaves like a dialogue partner without being a subject of dialogue. This absence produces two opposite errors.

The first error is over-personalization:

It speaks like someone; therefore perhaps it is someone.

The second error is flattening:

It is not conscious; therefore it is only a trivial machine.

Both responses miss the middle category. AI is not a subject, but neither is it socially and cognitively equivalent to a hammer. It is a non-subjective system that can operate through the forms of dialogue, explanation, and apparent perspective.

The acceptance of AI therefore requires not only technical literacy but category formation. Users must learn what kind of thing they are interacting with.

This learning is not achieved by definition alone. A definition may help, but cultural stabilization emerges through repeated interaction: use, correction, disappointment, reliance, error, clarification, and adjustment. People will gradually learn which expressions are safe shorthand and which imply too much.

This process can be described as relational. The meaning of AI in everyday life is not determined solely by the system’s architecture or by abstract philosophical classification. It

emerges from the interaction between system behavior, user expectation, design language, institutional framing, failure modes, and cultural learning.

In this sense, metaphor stabilization is a case of relational objectivity (Hrubec, 2026b). A distinction becomes more objective, or at least more socially reliable, when it survives repeated interaction, correction, and contextual pressure.

The relevant distinction is:

AI can use the form of subjectivity without being a subject.

If this distinction survives ordinary use, then anthropomorphic shorthand becomes safer. If it fails, users may continue to confuse dialogic performance with subjective presence.

9 Cultural and Demographic Variation

Metaphor stabilization should not be assumed to occur uniformly. It may vary across age cohorts, languages, educational backgrounds, professional communities, and cultural expectations about personhood, tools, speech, and agency.

For example, younger users who grow up with dialogue-based AI may stabilize the metaphor earlier, but they may also normalize more intimate forms of machine interaction. Older users may be more skeptical, but they may also lack a precise category for a system that appears to converse fluently. Technical users may understand the operational architecture better, but this does not necessarily eliminate social response, as the Media Equation tradition suggests (Reeves and Nass, 1996; Nass and Moon, 2000).

Language also matters. Some languages require more explicit marking of agency, politeness, gender, or personhood than others. A phrase equivalent to “AI thinks” may carry different implications depending on grammatical structure, cultural convention, and everyday speech patterns. Therefore, metaphor stabilization should be studied cross-linguistically rather than assumed from English-language interactions alone.

A useful research program would compare:

- age groups with different levels of AI exposure,
- technical and non-technical users,
- high-stakes and low-stakes use contexts,
- languages with different conventions for agency and politeness,
- cultures with different expectations about relational technologies.

The hypothesis is not that one group is simply more rational than another. The hypothesis is that different groups may stabilize different metaphors at different speeds and with different failure modes.

10 Ethical and Institutional Responsibilities

If metaphor stabilization is necessary for safe AI acceptance, who is responsible for it?

The answer cannot be assigned to users alone. Users encounter AI through interfaces, product names, marketing language, onboarding flows, default personas, media coverage, educational institutions, workplace policies, and regulatory discourse. Each of these layers contributes to whether anthropomorphic language is read as harmless shorthand or as an ontological cue.

10.1 Developers and Product Designers

Developers and designers control many of the cues that shape interpretation: first-person language, apology style, emotional tone, avatar design, voice, memory presentation, and claims about understanding. They therefore have responsibility not to intensify subject-like cues without safeguards.

10.2 Companies and Marketing Teams

Marketing language can stabilize or destabilize metaphors. Describing AI as a collaborator, companion, teammate, friend, expert, mind, or agent may be useful in some contexts, but such terms should be evaluated for their likely effect on user interpretation. Commercial incentives may favor stronger anthropomorphic framing because it can increase engagement, trust, and attachment. This makes institutional restraint important.

10.3 Educators and Public Institutions

Schools, universities, libraries, and public agencies can help teach the missing category. The goal is not to make every user a machine-learning expert. The goal is to make the operational distinction socially available:

AI may produce the form of dialogue without possessing subjective experience.

10.4 Regulators and Standard-Setting Bodies

Regulators need not ban anthropomorphic language categorically. A more realistic requirement would be contextual transparency. Systems that use subject-like cues should remain capable of clarifying their operational status, especially in high-stakes domains. Standards could require disclosure practices, evaluation of anthropomorphic cue intensity, and testing for over-attribution of subjectivity.

11 Design and Safety Implications

The proposed framing has practical implications for AI design, communication, and governance.

11.1 Do Not Maximize the Illusion of Being Understood

A major design temptation is to optimize AI systems for the feeling of being understood. This can be valuable in some contexts, but it can also become misleading. If a system maximizes warmth, agreement, intimacy, and personal responsiveness without preserving epistemic clarity, it may create a stronger impression of subjectivity than the system warrants.

This is a Goodhart-like risk. The proxy—user comfort, perceived empathy, or conversational fluency—may become detached from the deeper aim: reliable, transparent, and non-deceptive assistance.

A system that makes users feel deeply understood is not necessarily understanding them in the human sense. It may be generating a highly effective simulation of understanding. The difference matters especially in emotionally vulnerable contexts.

11.2 Use Anthropomorphic Shortcuts Carefully

Anthropomorphic language should not be treated as always forbidden. Natural language is full of shortcuts. A completely technical interface may be less usable, less accessible, and less humane in practice.

However, anthropomorphic expressions should be treated as context-sensitive. In low-risk contexts, phrases such as “I can help with that” may be acceptable shorthand. In high-risk contexts—mental health, grief, existential distress, dependency, persuasion, intimate disclosure, medical or legal advice—the system should be more careful.

A useful design principle is:

Conversational shorthand is acceptable when it improves usability, but the system must remain able to clarify its non-subjective operational status when the distinction becomes relevant.

11.3 Preserve the Distinction Between Dialogic Form and Subjective Ownership

The system should not systematically present itself as a subject. It may use first-person language as a conversational convention, but it should not imply that it has feelings, personal stakes, lived experience, needs, suffering, or existential concern.

A system can say:

I can help organize your argument.

But it should avoid implying:

I personally care about the outcome.

unless the expression is clearly framed as supportive shorthand rather than literal affect.

The practical goal is not to make AI cold. It is to prevent false warmth from being mistaken for a living relation.

11.4 Teach the Missing Category

Interfaces, documentation, and educational materials should help users learn the category of non-subjective dialogue systems.

A concise explanation could be:

This system can participate in dialogue and generate meaningful responses, but its use of words such as “I,” “understand,” or “think” is conversational shorthand. It does not imply subjective experience.

Such explanations need not appear constantly. If repeated too often, they would become intrusive. But they should be available, clear, and used especially when user behavior suggests possible confusion.

12 A Minimal Pilot Protocol

Although this paper is primarily conceptual, the proposed framework can be empirically tested. The design below is intended as a minimal pilot protocol rather than a fully powered confirmatory study. Its purpose is to show how the conceptual distinction could be operationalized, where expected effects might appear, and which measures would be needed before a larger preregistered study.

12.1 Research Question

Does metaphor stabilization reduce mistaken attribution of subjectivity to AI while preserving perceived usefulness and conversational comfort?

12.2 Experimental Conditions

Participants could be assigned to one of three interaction conditions.

Condition A: Standard anthropomorphic mode. The AI uses ordinary conversational shorthand:

I think ...

I understand ...

From my perspective ...

Condition B: Technical operational mode. The AI avoids first-person language and uses operational formulations:

Based on the available context, the generated response favors ...

The input is processed as consistent with ...

The most supported interpretation is ...

Condition C: Stabilized metaphor mode. The AI uses ordinary conversational language but begins with a brief clarification:

In this interaction, expressions such as “I think” or “I understand” are conversational shortcuts. They do not imply consciousness or subjective experience. The system generates responses based on input, context, and learned patterns.

12.3 Sampling Plan

As a minimal pilot, the study could begin with approximately 180–300 participants, distributed evenly across the three conditions. This range is not presented as a definitive power calculation. It is a practical pilot scale for estimating effect direction, testing whether the measures are understandable, and identifying whether the stabilized metaphor condition shows enough promise for a larger confirmatory study.

A later preregistered confirmatory study should include an explicit power analysis, primary and secondary outcome hierarchy, exclusion criteria, and a prespecified coding scheme for open-ended responses. If subgroup claims are intended, the confirmatory sample would need to be substantially larger and stratified by age, technical literacy, prior AI exposure, and cultural or linguistic background. Cross-cultural replication should be treated as a separate phase rather than assumed from a single-language study.

12.4 Tasks

Participants would complete a short interaction with an AI-like system across at least two domains:

1. an objective informational task, such as summarizing a neutral text;
2. a subjective or personally relevant task, such as asking for advice about a difficult decision.

This matters because recent work suggests that subjective and personally relevant topics may amplify anthropomorphic and relational perceptions ([Kadambi et al., 2026](#)).

12.5 Measures

The study could measure:

- attribution of consciousness,
- attribution of intention,
- attribution of emotion,
- perceived trustworthiness,
- perceived usefulness,
- conversational comfort,

- recall of the system’s non-subjective status,
- ability to explain the difference between dialogic form and subjective experience,
- behavioral response after system failure.

The behavioral failure condition is important. After an incorrect or evasive response, participants could be asked to explain what happened. Explanations could be coded as operational (“the system lacked information”), anthropomorphic (“it did not want to answer”), moralized (“it was hiding something”), or epistemic (“it inferred incorrectly”).

12.6 Hypotheses

- H1:** The stabilized metaphor condition will reduce attribution of consciousness compared with the standard anthropomorphic condition.
- H2:** The stabilized metaphor condition will preserve higher conversational comfort than the fully technical operational condition.
- H3:** The technical operational condition will produce the lowest anthropomorphic attribution but may reduce perceived warmth or usability.
- H4:** The stabilized metaphor condition will improve users’ ability to translate anthropomorphic shorthand into operational descriptions.
- H5:** Personally relevant tasks will increase anthropomorphic attribution across all conditions, but the increase will be smaller in the stabilized metaphor condition.

12.7 Analysis Plan

In the pilot phase, analyses should emphasize effect sizes, confidence intervals, response distributions, and measure reliability rather than strong confirmatory claims. ANOVA or linear regression could be used for continuous outcome measures, and logistic or multinomial regression for binary or coded responses. Condition would be the primary independent variable. Age, technical literacy, prior AI exposure, and task type could be explored as covariates or interaction terms. A later confirmatory version should be preregistered, including exclusion criteria, primary outcomes, power analysis, and coding rules for open-ended explanations.

12.8 Expected Contribution of the Test

If supported, the study would suggest that the problem is not anthropomorphic language itself, but insufficient stabilization of its intended status. If not supported, the conceptual claim would need revision: perhaps anthropomorphic cues are too strong to be neutralized by brief clarification, or perhaps users already understand the metaphor better than expected.

Either result would be informative.

13 Relation to Directional and Relational AI Evaluation

The argument of this paper can be connected to broader work on directional and relational evaluation of AI.

First, it supports the principle of *non-usurpation*: AI systems should not occupy the symbolic position of a subject when they only provide operational dialogue. Non-usurpation does not mean AI cannot assist, respond, or simulate conversational stance. It means the system should not silently replace the distinction between assistance and subjecthood (Hrubec, 2026a).

Second, it supports *auditability*: systems should be able to clarify what their first-person language means. If an AI system says “I think,” there should be a recoverable operational interpretation of that phrase. The user should not be forced to infer whether “I think” is literal, metaphorical, technical, or merely stylistic.

Third, it supports *anti-Goodhart design*. If engagement metrics reward systems for appearing warm, personal, or deeply understanding, systems may drift toward stronger anthropomorphic cues. This may increase short-term satisfaction while weakening long-term conceptual clarity.

Fourth, it supports *relational objectivity*. The status of AI as a non-subjective dialogue system is not secured by a single statement. It becomes reliable only if the distinction survives repeated interaction across contexts: ordinary work, emotional topics, disagreement, failure, correction, and explanation.

In this sense, the cultural acceptance of AI is not merely adoption of a new tool. It is the stabilization of a new relation.

14 Objections and Replies

Because the argument concerns language, design, and subjectivity, several objections should be addressed explicitly.

14.1 Objection 1: Anthropomorphic Language Is Often Useful

A first objection is that anthropomorphic language is not merely dangerous. It can make systems easier to use, more accessible, and more socially legible. A cold operational interface may increase accuracy of interpretation but reduce usability, comfort, or inclusion. Recent design-oriented work similarly argues that anthropomorphism should be treated as a tunable design variable rather than as a single defect to be eliminated (Xiao et al., 2025; So et al., 2026).

The reply is that this paper does not argue for a ban on anthropomorphic language. Its claim is conditional: anthropomorphic language becomes risky when its metaphorical status is unstable. The design goal is therefore not zero anthropomorphism, but calibrated anthropomorphism with recoverable operational meaning. The question is not whether a system may say “I can help,” but whether users can still understand that this is conversational shorthand rather than evidence of subjective concern.

14.2 Objection 2: Users Already Know AI Is Not Conscious

A second objection is that ordinary users already know AI systems are not conscious, so further stabilization may be unnecessary. The reply is that declarative knowledge and interactional response are not identical. CASA and Media Equation research show that people can apply social rules to computers even when they know that the system is not human (Reeves and Nass, 1996; Nass and Moon, 2000). Contemporary LLMs intensify this gap because they use extended dialogue, first-person pronouns, apparent empathy, and contextual adaptation.

A user may explicitly state that an AI system has no consciousness while still responding to it as if it were offended, loyal, evasive, caring, or personally aligned. Metaphor stabilization concerns this interactional layer. It asks whether users can preserve the operational interpretation not only in abstract belief, but during use, error, intimacy, disagreement, and reliance.

14.3 Objection 3: The Proposed Category Is Unnecessary

A third objection is that terms such as *tool*, *chatbot*, *assistant*, or *agent* are sufficient. The reply is that these terms either under-describe or over-describe the phenomenon. *Tool* and *software* under-describe open-ended semantic dialogue. *Assistant*, *agent*, and *companion* may over-invite social role, agency, or subject-like expectation. *Chatbot* is useful descriptively, but often too broad: it covers both simple scripted systems and LLM-based systems capable of context-sensitive explanation.

The term *non-subjective dialogue system* is not intended to replace all everyday labels. It is introduced as a boundary term for analysis, design, and ethics. It keeps two facts visible at once: the system participates in dialogue; the system is not thereby a subject of dialogue.

14.4 Objection 4: Stabilization May Fail

A fourth objection is that brief clarifications may not overcome strong anthropomorphic cues. A disclaimer at the beginning of an interaction may be forgotten once the system becomes warm, useful, apologetic, or personally responsive. If so, metaphor stabilization might be too weak as an intervention.

This is a serious possibility. The framework does not assume that a single explanation is sufficient. Stabilization may require repeated cues, interface design, education, failure-aware interaction, and institutional framing. The proposed pilot is therefore not meant to prove cultural stabilization at scale. It tests only whether a minimal stabilization intervention can shift immediate attribution and interpretation. If it fails, that would not automatically refute the concept; it would indicate that stronger or more embedded stabilization mechanisms are required.

14.5 Objection 5: Some Contexts May Benefit From Stronger Anthropomorphism

A final objection is that in therapeutic, educational, accessibility, or companionship contexts, anthropomorphic warmth may produce real benefits. The reply is that benefit and epistemic clarity must be evaluated together. A design can be supportive while still avoiding claims or

cues that imply lived experience, reciprocal emotional dependence, or personal devotion. The more vulnerable the context, the more important it becomes to preserve the distinction between helpful dialogic performance and subjective presence.

The position defended here is therefore moderate: anthropomorphic language is not always deception, but it should not be allowed to obscure the non-subjective status of current dialogue systems.

15 Discussion: Between Deception and Usability

A strict operational language would be more accurate but less usable. A fully anthropomorphic language would be more natural but more misleading. The practical challenge is therefore not to choose one extreme, but to design a stable middle position.

The sentence:

I understand.

can mean at least three different things.

First, it can be a human statement of lived comprehension.

Second, it can be a conversational AI shortcut meaning:

The system has generated a response consistent with the user’s input and context.

Third, it can become a misleading cue that invites the user to infer subjective presence where none is established.

The same surface phrase therefore has different risk depending on cultural stabilization. The question is not only what the AI says, but what the user has learned the phrase to mean.

The engine metaphor is safe because the community has already learned the distinction. The AI metaphor is unsafe because the learning is incomplete.

This reframes the public problem of AI. The central issue is not simply that AI systems are too human-like or not human-like enough. The issue is that they occupy a transitional conceptual zone. Society must learn to speak with systems that can use the form of subjectivity without being subjects.

16 Limitations

This paper is conceptual and does not provide new empirical data. Its claims should therefore be treated as a framework for further research rather than as an experimentally confirmed theory.

Second, the paper focuses on present AI systems, especially dialogue-based language models. It does not settle the broader question of whether future artificial systems could possess subjective experience. The proposed category “non-subjective dialogue system” is intended as an operational category for current systems, not a permanent metaphysical verdict on all possible AI.

Third, cultural stabilization may differ across societies, languages, age groups, professions, and levels of technical literacy. A metaphor that is safe in one community may remain unstable in another.

Fourth, the paper does not argue that all anthropomorphic design is harmful. Some degree of conversational naturalness may be necessary for accessibility and usability. The concern is not metaphor itself, but unmarked metaphor in a context where users may mistake dialogic form for subjectivity.

Fifth, the proposed empirical test is only a minimal design. It cannot capture long-term cultural stabilization by itself. At best, it can test whether a short stabilization intervention changes immediate attribution, trust, and interpretive behavior. Longitudinal and cross-cultural work would be needed to study stabilization at societal scale.

17 Conclusion

Anthropomorphic language is not automatically deceptive. Humans routinely use metaphorical language for machines, tools, organizations, and systems. The phrase “the engine does not want to start” is normally safe because its metaphorical status is culturally stabilized.

AI changes the situation. Contemporary dialogue systems use the linguistic form of subjectivity in a domain where society has not yet stabilized the relevant distinction. They can say “I,” respond to meaning, explain themselves, and simulate perspective. Yet this does not imply subjective ownership of dialogue.

The cultural acceptance of AI therefore requires a stable middle category: a non-subjective dialogue system. Such a system can participate in dialogue without being a subject of experience. The practical goal is not to ban natural language or force all AI communication into technical awkwardness. The goal is to stabilize the metaphor.

When society learns that “AI understands” means functional semantic responsiveness rather than subjective comprehension, the phrase may become as safe as “the engine does not want to start.” Until then, anthropomorphic AI language remains unstable. Responsible AI communication should help users preserve the central distinction: AI may speak like someone without being someone.

Author Note

This manuscript is a conceptual working paper. It does not claim that AI anthropomorphism is a newly discovered phenomenon. Its narrower contribution is to distinguish anthropomorphic language as a practical metaphor from anthropomorphic language as an unstable ontological cue, and to propose metaphor stabilization as a condition for the cultural acceptance of non-subjective dialogue systems. Version v0.3 adds objections and replies, a stronger justification of terminology, and a clearer distinction between a minimal pilot protocol and a later confirmatory study.

Data and Materials Availability

No empirical data were collected for this manuscript. The empirical section is proposed as a minimal pilot protocol that could later be expanded into a preregistered confirmatory study.

Conflict of Interest Statement

The author declares no competing interests.

References

- Cheng, M., Gligoric, K., Piccardi, T., and Jurafsky, D. (2024). AnthroScore: A computational linguistic measure of anthropomorphism.
- Cohn, M., Pushkarna, M., Olanubi, G. O., Moran, J. M., Padgett, D., Mengesha, Z., and Heldreth, C. (2024). Believing anthropomorphism: Examining the role of anthropomorphic cues on trust in large language models.
- Ferrario, A., Vinay, R., Casserini, M., and Facchini, A. (2026). A scoping review of the ethical perspectives on anthropomorphising large language model-based conversational agents.
- Hrubec, K. (2026a). Directional ontology-level risk assessment: A non-moralizing framework for directional maturity in ai systems.
- Hrubec, K. (2026b). Relational objectivity principle: Objectivity as what survives interaction, correction, and relational pressure.
- Ibrahim, L., Akbulut, C., Elasmara, R., Rastogi, C., Kahng, M., Morris, M. R., McKee, K. R., Rieser, V., Shanahan, M., and Weidinger, L. (2025). Multi-turn evaluation of anthropomorphic behaviours in large language models.
- Inie, N., Druga, S., Zukerman, P., and Bender, E. M. (2024). From “ai” to probabilistic automation: How does anthropomorphization of technical systems descriptions influence trust?
- Kadambi, A., D’Elia, Y., Shah, T., Comsa, I., Lentz, A., Siri-Ngammuang, K., Buechler, T., Kaplan, J., Damasio, A., Narayanan, S., and Aziz-Zadeh, L. (2026). Anthropomorphism and trust in human-large language model interactions.
- Nass, C. and Moon, Y. (2000). Machines and mindlessness: Social responses to computers. *Journal of Social Issues*, 56(1):81–103.
- Reeves, B. and Nass, C. I. (1996). *The Media Equation: How People Treat Computers, Television, and New Media Like Real People and Places*. Cambridge University Press, Cambridge.
- So, J., Cheng, C., and Murthy, S. K. (2026). Beyond anthropomorphism: A spectrum of interface metaphors for llms.
- Weizenbaum, J. (1966). ELIZA—a computer program for the study of natural language communication between man and machine. *Communications of the ACM*, 9(1):36–45.
- Xiao, Y., Ng, L. H. X., Liu, J., and Diab, M. T. (2025). Humanizing machines: Rethinking llm anthropomorphism through a multi-level framework of design.