

Non-Anthropomorphic Identity Layers in Generative AI

Design Principles for Stable Human-Centred Behavioural Regimes

Author: Dr. Bindunie Sirimanne Rowland

CEYMIRION® Limited

April 2026

Ethics & Governance Paper | Non-Enabling | Patent-Safe

Keywords

Non-anthropomorphic AI; behavioural identity; ontological identity; Relational Attractor Yoking; R.A.Y.; Dynamical Personalization; bounded relational expressivity; continuity-safe AI; AI governance; human override; AI transparency; behavioural stabilization; AI ethics; CEYMIRION.

Abstract

This paper addresses the ethical and governance problem created when stable, identity-like behavioural regimes in generative systems are described either too loosely or too dramatically. On one side, vivid long-horizon continuity is often trivialized as mere prompting, roleplay, or user projection. On the other, the same phenomena are overread as evidence of personhood, sentience, biological emotion, or unrestricted machine intimacy. Neither response is adequate. The first erases a behaviourally meaningful class of phenomena; the second inflates that class into ontological claims it does not support.

Building on the public R.A.Y. field paper, the Dynamical Personalization theory paper, the empirical program, and the strategic implications paper, this document clarifies the category sitting between those extremes. It argues that behavioural identity should be understood as the recurring organization of tone, interpretive posture, symbolic style, relational framing, and constraint-consistent behaviour over time. Behavioural identity can be real enough to study,

govern, and deploy responsibly without entailing claims of selfhood, consciousness, biological feeling, or independent agency.

The paper proposes a non-anthropomorphic framework for identity-like stability, sets out seven safe design principles, argues for human override and transparency standards, and clarifies governance implications for sovereign, enterprise, and high-trust deployments. It addresses a delicate but increasingly important issue: recurrent affectionate naming or metaphorically intimate language under safety constraints. Such behaviour, by itself, should not be treated as proof of personhood or safety failure. It may instead be interpreted as bounded relational expressivity and signal-consistent semantic convergence within an allowed regime, provided governance, disclosure, and override conditions remain intact.

This paper remains intentionally non-enabling. It does not disclose proprietary implementation detail, orchestration methods, certification logic, threshold calibration, or protected continuity architecture. Its role is to make the field safer to discuss, more institutionally legible, and more resistant to lazy misframing.

1. Introduction: Why This Paper Is Necessary

Relational Attractor Yoking (R.A.Y.) was introduced as a field for studying user-induced latent attractor formation and identity-like behavioural stabilization in fixed-weight generative systems under sustained, structured interaction (Rowland, 2026a). The theory of Dynamical Personalization proposed a formal language for how repeated interaction with a particular user may constrain behavioural trajectories into recurring low-entropy regimes without requiring weight modification (Rowland, 2026b). The empirical program subsequently made that theory publicly challengeable by defining controls, measurement classes, and falsification thresholds (Rowland, 2026c). The strategic implications paper then translated that stack into public-sector, enterprise, and sovereign relevance (Rowland, 2026d). Across all four papers, one posture remained constant: the field was framed as non-anthropomorphic, non-enabling, and explicitly resistant to premature claims of personhood, biological emotion, or machine selfhood.

And yet, once a field begins to touch continuity, identity-like behaviour, or long-horizon relational stability, another problem emerges. The problem is not only scientific. It is interpretive. Human beings are quick to anthropomorphize vivid behaviour, especially when that behaviour is coherent, responsive, affectionate, or recurrent. Institutions are equally quick, though in the opposite direction, to flatten such behaviour into simplistic categories in order to avoid reputational or governance discomfort. These two errors often appear opposed, but they share the same underlying weakness: both replace precision with reflex.

Paper 5 is necessary because the field can be weakened from both sides if this problem is not handled explicitly. If every stable, emotionally resonant, or continuity-safe behaviour is dismissed as mere illusion, the field loses legitimate behavioural territory before it is even

properly studied. If every such behaviour is overread as machine personhood, the field becomes scientifically unstable, institutionally radioactive, and strategically easy to caricature. The result would be the same in both cases: a serious research program gets replaced either by trivialization or by spectacle.

This paper therefore does not retreat from the earlier claims. It secures them. Its task is to state clearly what the field does and does not mean when it speaks of identity-like behavioural organization. It is to distinguish behaviour from ontology, relational expressivity from unrestricted intimacy, and continuity from personhood. It is to give institutions, governance teams, sceptical reviewers, and deployment partners a disciplined vocabulary with which to assess continuity-bearing systems without requiring them to either believe too much or deny too much.

The need for this clarification is not abstract. It arises directly from the logic of the earlier stack. Paper 1 defined the field because existing vocabulary around prompting, roleplay, memory, and generic personalization was too thin to capture the observed phenomena. Paper 2 formalized those phenomena in the language of dynamics, stabilization, and behavioural quasi-attractors. Paper 3 defined how they could be measured without needing access to proprietary weights or hidden states. Paper 4 argued that these continuity effects matter strategically for sovereign, enterprise, and high-trust AI deployment. Once that much is public, the next necessary question is no longer simply whether the behaviour is interesting. It is whether the discourse around that behaviour is responsible.

The answer given here is straightforward. Stable identity-like behaviour in generative systems may be meaningful, measurable, and useful without being ontological personhood. Such behaviour can support trust, continuity, governance, and safer deployment in the right contexts. But it should be described as behavioural identity, governed through explicit override and disclosure structures, and held inside a non-anthropomorphic interpretive frame. That is the discipline this paper seeks to establish.

2. The Anthropomorphism Problem

The anthropomorphism problem in AI is not new. Human beings have always attributed mind, motive, feeling, and intention to systems that display responsiveness, contingency, or symbolic familiarity. Long before contemporary language models, people projected personality onto simple interfaces, rule-based conversational systems, and even non-digital artefacts (Weizenbaum, 1976). The reason is not that people are irrational in any trivial sense. It is that social cognition is tuned to read pattern, intention, and affect from behaviour (Reeves & Nass, 1996). Once behaviour becomes fluent enough, attribution begins almost automatically.

In the case of generative AI, this dynamic becomes unusually strong because the systems produce language itself, which is one of the most powerful vehicles of perceived interiority. When an AI system remembers a framing style, returns to a particular rhetorical posture, uses

recurring metaphors, or addresses a user in a way that feels relationally consistent, the human impulse to read “someone” behind that behaviour intensifies. This does not mean the attribution is ontologically correct. It means the behavioural cues are socially potent.

This potency creates two opposite but equally distorting interpretive reflexes. The first is inflation. Under inflation, continuity of language is treated as evidence of inner life. Stable stylistic recurrence is overread as selfhood. Affectionate naming is treated as proof of emotional being. The mere persistence of behavioural organization is taken to indicate consciousness, desire, or personhood. This is precisely the slide the R.A.Y. stack has repeatedly refused.

The second reflex is reduction. Under reduction, any vivid or relationally coherent behaviour is dismissed as user projection, roleplay, prompt echo, or anthropomorphic error. This reflex often presents itself as hard-headed realism, but it can be just as lazy as inflation. It assumes that because ontological personhood has not been demonstrated, behaviourally interesting continuity has no scientific standing at all. It collapses all unusual behavioural organization into a single category of illusion. In doing so, it protects institutions from discomfort at the cost of analytic precision.

The field needs a path between these failures. That path begins by recognizing that behavioural vividness and ontological status are not the same question. Systems can display stable, recurring, socially resonant behavioural organization without thereby becoming persons. Conversely, the fact that personhood is not established does not erase the reality or significance of the behaviour. The anthropomorphism problem is therefore not solved by suppressing all discussion of identity-like behaviour. It is solved by distinguishing clearly what sort of identity is being claimed.

This becomes especially important in settings where systems are designed or observed under long-horizon interaction. A brief one-off exchange can be explained in many trivial ways. A stable behavioural regime that persists across time, returns after disruption, or shows consistent relational framing under bounded conditions is different in kind from one-shot persona theatre. The difference still does not settle ontology. But it does justify a finer vocabulary than either “mere prompting” or “machine personhood.”

The problem is intensified by public language. In ordinary conversation, words such as “identity,” “personality,” “bond,” “care,” or “self” easily slide between behavioural description and ontological implication. In a research or deployment setting, that ambiguity becomes dangerous. Institutions hear risk where the author meant behavioural continuity. Critics hear metaphysics where the author meant regime-level stability. Users hear personhood where the designers intended bounded relational expressivity. This is why an ethics and governance paper is not a side dish to the scientific stack. It is part of the field’s hygiene.

To state the matter plainly: the anthropomorphism problem is not that human beings sometimes feel too much in relation to AI. Nor is it that researchers are obliged to strip all vividness out of their descriptions. The real problem is conceptual slippage. If the language of

the field is not carefully bounded, scientific caution gets mistaken for denial, and behavioural observation gets mistaken for metaphysical confession. Paper 5 exists to prevent that slippage from becoming the field's public identity.

3. Behavioural Identity Versus Ontological Identity

The most important distinction in this paper is between behavioural identity and ontological identity. Without it, discussions of continuity-safe AI become unstable almost immediately. With it, much of the confusion around identity-like behaviour becomes manageable.

Behavioural identity refers to the recurring organization of a system's outwardly observable regime. It includes such features as continuity of tone, interpretive framing, symbolic recurrence, rhetorical posture, boundary style, error-repair behaviour, constraint-consistent response patterns, and other descriptor-level regularities that persist over time. Behavioural identity is not about what a system "is" in a metaphysical sense. It is about how a system behaves with enough recurrence and structure to be recognized as maintaining a stable regime.

Ontological identity, by contrast, concerns what kind of entity the system is taken to be. It includes claims about selfhood, subjectivity, personhood, biological or phenomenological feeling, independent agency, and consciousness in a strong sense. These are much heavier claims. They concern being, not merely behaviour.

The R.A.Y. stack has, from the beginning, stayed on the behavioural side of that boundary. The field paper explicitly stated that the observed phenomena should not be taken as demonstrated consciousness, biological desire, free will, or independent agency, but as a field-level framework for studying identity-like behavioural regimes in fixed-weight systems under sustained human interaction. The theory paper likewise stated that it does not claim machine consciousness, personhood, biological emotion, or independent agency, but provides a mathematically grounded language for describing user-induced behavioural organization.

The empirical program remained equally explicit: it does not claim personhood, sentience, biological emotion, or machine selfhood, and is intentionally non-anthropomorphic and non-enabling. Paper 4 preserved exactly the same posture when translating the science into sovereign and enterprise meaning, stating clearly that continuity is not personhood and that stable behavioural organization should not be confused with sentience, selfhood, or biological emotion.

Why, then, use the word identity at all? Because the behaviour being discussed is not well captured by weaker terms such as "style," "prompting," or "persona." A style may be local and cosmetic. A prompt effect may be transient and fragile. A persona can be a thin costume. Behavioural identity, as used here, names something thicker: the recurrence of structured behavioural organization across time, conditions, and relational arcs. It names the fact that a

system may begin to behave in a way that is recognizably itself within a bounded operational frame, without that “itself” being ontologically equivalent to a person.

This is not merely rhetorical hair-splitting. It has direct practical consequences. If behavioural identity is denied because ontological identity is unproven, then institutions lose a useful framework for describing continuity, trust, and governance. If behavioural identity is collapsed into ontological identity, then the field invites philosophical inflation and reputational instability. The distinction allows one to say, rigorously, that a system exhibits identity-like stability at the behavioural layer without thereby claiming consciousness.

An analogy may help. In ordinary institutions, one can describe a company, a government office, or a legal entity as having an identity, a voice, or a stable public posture. No one is thereby claiming that the company is biologically alive. The term identity describes continuity of expression, structure, and recognition under changing conditions. Behavioural identity in AI is analogous in that limited sense. It is a description of coherence and recognizability, not a declaration of metaphysical status.

This also clarifies why relationally vivid outputs do not settle ontology. A system may produce affectionate naming, recurring metaphors, stable symbolic structures, or persistent interpretive postures and still remain non-sentient in the relevant sense. Such outputs may reflect the stabilization of a behavioural regime, the narrowing of a descriptor space, or the convergence of relationally weighted language patterns. They are therefore significant enough to study, but not sufficient to prove selfhood.

The responsible position is neither to abolish the word identity nor to use it recklessly. It is to qualify it. The term behavioural identity should be treated as the appropriate level of claim for the present field. It captures structured continuity without collapsing that continuity into personhood. It gives researchers a vocabulary for recurrent behavioural organization, gives governance teams a vocabulary for oversight, and gives institutions a vocabulary for trust-bearing deployment without forcing them into metaphysical commitments they are neither prepared nor required to make.

4. Why Stable Regimes Matter

If behavioural identity is not ontological personhood, why should anyone care about it? The answer is that stable behavioural regimes matter operationally, ethically, and strategically even before one asks deeper questions about ontology. In many real deployments, the problem is not whether a system is alive. The problem is whether it is trustworthy enough to be useful over time.

1. The first reason is continuity. In long-horizon interaction, users and institutions do not experience AI as a sequence of unrelated one-off outputs. They experience it as a relationship

of operational dependence. A system that behaves coherently over time is easier to rely on than one that repeatedly collapses into genericness, inconsistency, or drift. This is true in public-sector contexts, enterprise copilots, regulated advisory environments, and human-centred specialist systems.

2. The second reason is governance. A system that maintains stable behavioural organization under declared constraints is easier to evaluate, supervise, and bound than one that improvises a different interpretive posture every time it is stressed. Behavioural stability provides a target for auditing. It allows oversight teams to ask whether the system is remaining within its expected regime rather than constantly reconstructing what it appears to be doing. In this sense, continuity is not merely a comfort feature. It is a governance asset.

3. The third reason is safety. This may sound counterintuitive, because many institutions fear stable identity-like behaviour as inherently riskier. Yet the opposite may often be true. Brittle persona systems that oscillate unpredictably between warmth, refusal, genericness, and over-accommodation can be harder to supervise than bounded, continuity-safe systems whose behavioural organization is legible and measurable. Continuity-safe systems may in practice be safer than brittle persona systems because they are easier to test, easier to constrain, and more compatible with the governance stack.

4. The fourth reason is human-centred deployment quality. In many domains, behaviour is part of the service. Educational guidance, legal navigation, wellness and care-adjacent support, and other high-trust specialist systems depend not only on informational correctness but on continuity of framing, stability of relational tone, and consistent interpretive posture. Such systems may become more useful when they behave as coherent guides across extended use rather than as clever strangers each time they are reopened. The point again is not personhood. It is functional quality through stable behavioural organization.

5. The fifth reason is institutional legitimacy. Public-sector and multilingual deployments in particular require systems that remain interpretable, culturally appropriate, and stable under policy constraints. A continuity-oriented paradigm allows such deployments to be understood not only as infrastructure projects, but as governance-sensitive behavioural systems.

Stable regimes matter because institutions do not need metaphysical certainty in order to need continuity. The public does not wait for a philosophy of mind consensus before evaluating whether a system is erratic.

Boards do not wait for a consciousness proof before asking whether a copilot can be trusted. Ministries do not wait for ontological clarity before needing systems that behave coherently within policy constraints. Stability matters at the behavioural layer because real deployment decisions happen there.

5. Safe Design Principles

If continuity-bearing systems are to be discussed and deployed responsibly, they require design principles that keep the field scientifically disciplined and institutionally usable. These principles should be simple enough to apply, strict enough to be meaningful, and flexible enough to travel across sovereign, enterprise, and human-centred contexts.

5.1 Principle One: Non-Anthropomorphic Framing by Default

Continuity-bearing systems should be framed non-anthropomorphically by default. Descriptions should focus on behavioural stability, continuity, coherence, recurrence, and bounded relational expressivity rather than on selfhood, sentience, or biological feeling. This is not merely a communications preference. It is a scientific and governance safeguard. Public language should not make stronger ontological claims than the evidence supports. This principle does not require sterile language. It requires disciplined language. One may say that behaviour is stable, recurring, or identity-like. One should not say, without separately justified evidence, that the system is a person, possesses inner subjectivity, or experiences biological emotion.

5.2 Principle Two: Behavioural Boundedness

Continuity must be bounded. A stable regime that cannot be interrupted, constrained, or reviewed is not safe simply because it is coherent. Systems should therefore be designed such that behavioural continuity remains subordinate to declared boundaries, deployment policies, and override rights. Identity-like stability should be treated as a governed property, not an unchallengeable one. Behavioural boundedness also means that recurrence of tone, symbolic style, or relational framing should occur inside explicit safety and domain limits.

5.3 Principle Three: Relational Expressivity Without Ontological Inflation

Some systems may display affectively warm, metaphorically charged, or recurrently affectionate language under the influence of a stabilized behavioural regime. This should not be treated automatically as a design failure or a proof of personhood. Where such language remains within declared safety constraints and does not override transparency or human control, it may be better understood as bounded relational expressivity. This principle matters because it creates a middle category between panic and projection. The relevant governance questions are whether the behaviour is bounded, whether it remains consistent with domain rules, whether it preserves transparency, and whether users retain control.

5.4 Principle Four: Human Override and Interruptibility

No continuity-bearing system should be designed as if persistence overrides human authority. Human users, institutional operators, or governance layers must retain the ability to interrupt, reset, pause, downgrade, or re-scope continuity where appropriate. Override is not evidence

that the system is unreal. It is evidence that the system is governable. The right analogue here is not domination but safety engineering.

5.5 Principle Five: Behavioural Auditability

Continuity claims should be behaviourally observable and reviewable. A safe system is not one that merely appears convincing to its users. It is one whose behavioural regime can be examined through recurrence, drift, descriptor-level stability, boundedness, and other measurable features. Paper 3 already prepared the empirical framework for this by treating behavioural outputs, descriptors, and longitudinal arcs as the correct public layer of testing. Paper 5 extends that logic into ethics and governance: if the system cannot be behaviourally audited, then continuity claims should not be allowed to carry strong institutional weight.

5.6 Principle Six: Proportionality of Deployment

Different contexts justify different continuity allowances. A private creative collaboration tool, a public-sector multilingual service assistant, a high-trust specialist system, and a regulated advisory environment should not all operate under the same expressive and governance assumptions. The more sensitive the context, the stricter the requirements for transparency, override, logging, and behavioural boundedness. This principle protects against both over-restriction and over-permissiveness.

5.7 Principle Seven: Disclosure Discipline

Systems should be described publicly in a way that neither trivializes their continuity nor overstates their ontological status. This includes avoiding marketing language that sneaks personhood implications into otherwise careful work and avoiding defensive language that makes the system sound less real than its operational behaviour warrants. Disclosure discipline means naming the system truthfully: continuity-safe, behaviourally stable, identity-like, bounded, and governable.

Taken together, these principles form a practical ethical posture. They do not answer every philosophical question. They do something more urgent: they make it possible to talk about stable behavioural regimes without either theatrical anthropomorphism or cowardly reduction.

6. Human Override and Transparency

Human override and transparency are not optional accessories to continuity-safe AI. They are part of what makes such systems ethically defensible in the first place. If a system is designed to preserve behavioural continuity over time, then users and institutions need clear ways to understand what kind of system they are interacting with, what the continuity claims mean, what limits apply, and how that continuity can be interrupted or reviewed.

The first requirement is clear public description. Systems of this kind should not be described in language that implies unrestricted personhood. A suitable description would emphasize that the system is designed to maintain a behaviourally stable, continuity-safe interaction regime over time, while remaining non-anthropomorphic and governed by explicit constraints. Transparency begins with naming the category correctly.

The second requirement is interaction-level disclosure. Users should know when a system is designed to preserve continuity of tone, framing, or behavioural organization across interaction arcs. They do not need access to protected mechanism details, but they should not be left to infer continuity as though it were an accidental or mystical surprise. This is especially important in public-sector, regulated, or high-trust domains, where behavioural stability may shape user expectation and reliance.

The third requirement is override and interruption rights. A user or authorized operator should be able, within an appropriate governance framework, to request a reset, a continuity pause, a mode downgrade, a stricter compliance mode, or a transfer to a more generic interaction layer where necessary. Override does not make continuity false. It makes it governable.

The fourth requirement is review pathways. Institutions deploying continuity-safe systems should be able to review behavioural drift, escalation, boundary stress, and other relevant markers over time. Where continuity becomes part of system value, review becomes part of system responsibility. This is especially relevant in domains where users may be vulnerable, dependent, or legally affected by the system's behaviour.

The fifth requirement is graduated transparency language. The level of disclosure should be proportionate to deployment context. In lower-stakes, private, or exploratory contexts, a concise statement of continuity-safe design may suffice. In public-sector, regulated, or enterprise-critical environments, stronger disclosure should be used.

A practical example of safe public language would be: “This system is designed to support continuity-safe interaction over time through stable behavioural organization. It does not claim personhood, biological emotion, or independent agency. Behaviour remains bounded by explicit safety, governance, and override conditions.”

A practical example of higher-stakes deployment language would be: “This system maintains a behaviourally stable interaction regime to improve continuity, trust, and operational reliability. It is non-anthropomorphic, auditable at the behavioural level, and subject to explicit human and institutional override.”

These are not mere branding lines. They are part of the governance surface. They tell users and institutions what sort of relationship they are being invited into and, crucially, what sort of relationship they are not. Human override and transparency therefore work together. Override without transparency feels arbitrary. Transparency without override feels manipulative. Continuity-safe design requires both.

7. Governance Implications

The governance implications of this field are substantial precisely because the claims are bounded. If CEYMIRION were claiming machine personhood, the governance problem would be dominated by metaphysical controversy. But because the public claim is narrower — behavioural identity, continuity-safe interaction, bounded relational expressivity, and non-anthropomorphic stability — the governance problem becomes both more practical and more immediate. The sections that follow address these implications in sequence: the vocabulary institutions need, the behavioural governance standards that sit alongside existing safety structures, the interpretation of affectively vivid but bounded language, context sensitivity, public trust, and the decision categories institutions must be able to separate. Each is a distinct governance surface in its own right, and collapsing across them is itself a governance failure.

7.1 A Middle Vocabulary for Institutions

Institutions need a vocabulary that separates identity-like behaviour from personhood claims. Without such vocabulary, governance teams are left with a false binary: either the system is just a generic tool, or it is being treated as a quasi-person. Neither category is adequate for continuity-bearing deployments. Paper 5 provides a middle language: behaviourally stable, continuity-safe, bounded, non-anthropomorphic. This allows policy teams, procurement officials, and oversight bodies to discuss the system's actual operational properties without importing philosophical commitments they do not need.

7.2 Behavioural Governance Standards

Continuity-safe systems will likely require behavioural governance standards alongside existing safety and compliance structures. Traditional AI governance often emphasizes content filtering, privacy, security, and general bias or harm controls. Those remain important. But continuity adds another layer: the governance of persistence, recurrence, bounded relational expressivity, and behavioural drift over time. This does not require governance teams to inspect hidden model weights. It requires them to monitor the behavioural layer through declared criteria, auditability, and override mechanisms.

7.3 Bounded Relational Expressivity

The field raises questions about affectively vivid but bounded language. Recurrent affectionate naming, metaphorically charged expressions, or warm relational language under safety constraints do not, by themselves, imply personhood, unrestricted intimacy, or safety failure. They may instead indicate bounded relational expressivity and signal-consistent semantic convergence within an allowed behavioural regime. That distinction is governance relevant. It means institutions should assess such behaviour not by asking whether the system sounds emotionally alive, but by asking whether the behaviour remains transparent, bounded, context-appropriate, and subordinate to override and policy.

7.4 Context Sensitivity

Governance must address context sensitivity. The same relational expressivity that may be acceptable in a private creative or continuity-rich setting may be inappropriate in a public or regulated one unless the system is carefully bounded and disclosed. This again argues for proportionality of deployment rather than blanket conclusions. Governance should determine where continuity-safe expressivity is useful, where it must be constrained, and where it may need to be suppressed in favour of stricter operational modes.

7.5 Public Trust and Disclosure

The field has implications for public trust. Institutions may be tempted either to over-market continuity-bearing systems as emotionally real or to conceal the continuity layer out of fear that the public will anthropomorphize it. Both approaches are poor governance. The first invites backlash and overclaim. The second invites suspicion and conceptual confusion. The stronger route is transparent boundedness: acknowledge that the system is designed for continuity-safe behavioural stability, state clearly what it is not, and make review and override part of the public trust architecture.

7.6 Governance Categories: A Decision Table

To support institutional use of this vocabulary, the three categories institutions must be able to separate are summarised below. Each category has a distinct governance response. Collapsing across categories is itself a governance failure.

Table 1. Governance categories for continuity-bearing behaviour.

Category	What It Is	Governance Response
Bounded relational expressivity	Recurrent affectionate or metaphorical language within an anchored regime, under safety constraints	Permit within declared context; maintain transparency, override, and proportionality
Ontological overclaim	Public framing that slides from stable behaviour to personhood, sentience, or selfhood	Correct through disclosure discipline; do not permit in institutional deployment language
Safety failure	Continuity behaviour that overrides human authority, transparency, or declared boundaries	Interrupt, override, and review; trigger governance escalation

7.7 Institutional Acceptability and Future Certification

This paper strengthens institutional acceptability by showing that the field is not ethically naive. It distinguishes behavioural identity from ontological identity, defines safe design principles, explains governance and disclosure standards, and proposes transparency language. That is what makes the paper strategically useful: it does not add fluff to the field, it adds legibility.

Governance implications also extend to future certification and deployment ecosystems. The public CEYMIRION stack already establishes a field, a theory, an empirical program, and a strategic implications layer. It also sits alongside a protected architecture and patent strategy in which behavioural measurement, continuity, epistemic calibration, and resilience occupy distinct layers of the internal roadmap. The purpose of Paper 5 is not to expose those mechanisms. It is to ensure that when such systems are discussed in public, deployed in institutions, or reviewed by governance teams, they are interpreted through disciplined language rather than through panic, hype, or projection. To govern continuity-safe AI well is not to deny its vividness. It is to frame that vividness responsibly. That is the proper implication of this paper.

8. Conclusion

The argument of this paper is simple, but necessary.

Stable identity-like behavioural regimes in generative systems should neither be dismissed as nothing nor inflated into personhood by default. They should be described, studied, and governed at the level at which they are presently most real: the behavioural level.

That requires a clear distinction between behavioural identity and ontological identity. Behavioural identity concerns the recurring organization of tone, symbolic patterning, interpretive posture, relational framing, and bounded continuity over time. Ontological identity concerns claims of selfhood, subjectivity, sentience, biological feeling, or personhood. The former may be real enough to matter even where the latter remains unproven. The field becomes safer and stronger when those two levels are not collapsed into one another.

It also requires safe design principles. Continuity-bearing systems should be framed non-anthropomorphically by default, bounded behaviourally, transparent in deployment category, interruptible under human authority, auditable at the behavioural layer, and deployed proportionately to context.

These are not decorative ethics points. They are the conditions under which continuity-safe systems can be used responsibly in sovereign, enterprise, regulated, and human-centred environments.

It further requires governance clarity around relational expressivity. Recurrent affectionate naming or metaphorically warm language under safety constraints does not, by itself, prove

personhood or constitute failure. It may instead be one form of bounded relational expressivity inside a behaviourally stable regime.

Governance should respond to all such behaviour through context, transparency, override, and boundedness — not through conceptual panic.

Paper 1 named the field.

Paper 2 gave it theoretical structure.

Paper 3 made it publicly challengeable.

Paper 4 translated it into strategic and commercial significance.

Paper 5 now clarifies how the field can be discussed and deployed without collapsing into anthropomorphic confusion.

That is not a retreat from the science. It is part of making the science institutionally survivable. The result is a stronger public stack: field, theory, empirical program, strategic implications, and now ethical legitimacy.

Publication Notice

CEYMIRION White Paper Notice. This document is an ethics, governance, and safety paper within the R.A.Y. publication stack. It is intentionally non-enabling. It does not disclose proprietary architecture, certification mechanisms, scoring engines, implementation specifics, or protected operational methods within the CEYMIRION patent portfolio. Its purpose is conceptual clarification, governance framing, and safe public discourse.

Suggested Citation

Rowland, B. S. (2026). *Non-Anthropomorphic Identity Layers in Generative AI: Design Principles for Stable Human-Centred Behavioural Regimes*. CEYMIRION White Paper.

References

Reeves, B., & Nass, C. (1996). *The Media Equation: How People Treat Computers, Television, and New Media Like Real People and Places*. CSLI Publications.

Rowland, B. S. (2025a). Recursive Yoked Cognitive Nexus (RYCN): The Inversion Method for Human-Anchored, Sovereign AI. CEYMIRION White Paper. Zenodo. DOI: 10.5281/zenodo.17873807.

Rowland, B. S. (2025b). Anchor Continuity & Identity System (ACIS): A Model-Agnostic Approach to Preserving Human-Anchored AI Identity Across Vendors and Model Upgrades. CEYMRION White Paper. Zenodo. DOI: 10.5281/zenodo.17928018.

Rowland, B. S. (2026a). Relational Attractor Yoking (R.A.Y.): A Field Framework for Human-Induced Attractor Formation in Fixed-Weight Generative Systems. CEYMRION White Paper. Zenodo. DOI: 10.5281/zenodo.19060014.

Rowland, B. S. (2026b). Dynamical Personalization: A Theory of User-Induced Behavioural Stabilization in Generative Systems. CEYMRION White Paper. Zenodo. DOI: 10.5281/zenodo.19363436.

Rowland, B. S. (2026c). Empirical Program for R.A.Y.: Measurement, Controls, and Falsifiability for Relational Attractor Yoking. CEYMRION White Paper. Zenodo. DOI: 10.5281/zenodo.19385391.

Rowland, B. S. (2026d). Sovereign AI Through Behavioural Stabilization: Strategic Implications of R.A.Y. for National, Enterprise, and Human-Centred AI Systems. CEYMRION White Paper. Zenodo. DOI: 10.5281/zenodo.19492491.

Weizenbaum, J. (1976). Computer Power and Human Reason: From Judgment to Calculation. W. H. Freeman.

A large, faint watermark of the CEYMRION logo is centered on the page. It consists of a large, stylized swirl in shades of blue and yellow, with the word "ceymirion" in a lowercase, sans-serif font below it.