

When Generated Words Become Clinical Acts: An Austinian Account of Conversational Artificial Intelligence

Nicola Luigi Bragazzi, Halil Ibrahim Ceylan, Sergio Garbarino

Submitted to: JMIR AI
on: August 05, 2026

Disclaimer: © The authors. All rights reserved. This is a privileged document currently under peer-review/community review. Authors have provided JMIR Publications with an exclusive license to publish this preprint on its website for review purposes only. While the final peer-reviewed paper may be licensed under a CC BY license on publication, at this stage authors and publisher expressly prohibit redistribution of this draft paper other than for review purposes.

Table of Contents

Original Manuscript.....	4
---------------------------------	----------

Preprint
JMIR Publications

When Generated Words Become Clinical Acts: An Austinian Account of Conversational Artificial Intelligence

Nicola Luigi Bragazzi¹ MD, PhD, MPH; Halil Ibrahim Ceylan² PhD; Sergio Garbarino³ MD, PhD

¹ York University Toronto CA

² Atatürk University Erzurum TR

³ University of Genoa Genoa IT

Corresponding Author:

Nicola Luigi Bragazzi MD, PhD, MPH

York University
Keele Street 4700
Toronto
CA

Abstract

Generative and conversational artificial intelligence (AI) are commonly evaluated as systems that produce information. In health care, however, language does not merely represent states of affairs: it warns, reassures, recommends, documents, elicits disclosure, allocates attention, and can precipitate action. Drawing on J. L. Austin's theory of speech acts, this Viewpoint argues that the central philosophical and governance problem of conversational AI is not whether a large language model is a speaker in the full human sense. The more urgent problem is that AI-generated language can achieve perlocutionary efficacy without bearing the illocutionary commitments that ordinarily make speakers answerable for assertions, promises, warnings, and advice. This asymmetry is intensified by conversational interfaces, institutional branding, clinical workflow integration, and users' uptake of generated text as authoritative. The paper develops the concept of sociotechnical felicity: the alignment among linguistic force, evidence, delegated authority, user understanding, downstream consequences, and mechanisms of repair. It also proposes an Austinian audit for clinical conversational AI that extends evaluation beyond factual accuracy and surface empathy to force calibration, authorization integrity, uptake, perlocutionary safety, and contestability. This approach neither anthropomorphizes models nor treats them as inert tools. It locates intentional and moral responsibility in people and institutions while recognizing the model's causal contribution to open-ended communicative acts. The practical implication is direct: clinical conversational AI should be governed not only as a generator of content, but as infrastructure for actions performed through words.

(JMIR Preprints 05/08/2026:108858)

DOI: <https://doi.org/10.2196/preprints.108858>

Preprint Settings

1) Would you like to publish your submitted manuscript as preprint?

✓ **Yes – Make my manuscript publicly available as a preprint immediately.**

No – I do not wish to publish a preprint.

2) If accepted for publication in a JMIR journal, would you like the PDF to be visible to the public?

✓ **Yes, please make my accepted manuscript PDF available to anyone at any time (Recommended).**

No. Please do not make my accepted manuscript PDF available to anyone. I understand that if I later pay to participate in <https://preprints.jmir.org/preprint/108858>

Original Manuscript

Viewpoint

When Generated Words Become Clinical Acts: An Austinian Account of Conversational Artificial Intelligence

Nicola Luigi Bragazzi^{1,2}; Halil Ibrahim Ceylan³; Sergio Garbarino⁴

1 Laboratory for Industrial and Applied Mathematics (LIAM), Department of Mathematics and Statistics, York University, Toronto, ON, Canada.

2 Department of Computer Science, Data Science, and Information Technology, Faculty of Natural and Applied Sciences, Sol Plaatje University, Kimberley, South Africa.

3 Physical Education and Sports Teaching Department, Faculty of Sports Sciences, Atatürk University, 25240 Erzurum, Türkiye.

4 Department of Neuroscience, Rehabilitation, Ophthalmology, Genetics and Maternal/Child Sciences (DINOEMI), University of Genoa, 16132 Genoa, Italy.

Corresponding author: Nicola Luigi Bragazzi

Abstract

Generative and conversational artificial intelligence (AI) are commonly evaluated as systems that produce information. In health care, however, language does not merely represent states of affairs: it warns, reassures, recommends, documents, elicits disclosure, allocates attention, and can precipitate action. Drawing on J. L. Austin's theory of speech acts, this Viewpoint argues that the central philosophical and governance problem of conversational AI is not whether a large language model is a speaker in the full human sense. The more urgent problem is that AI-generated language can achieve perlocutionary efficacy without bearing the illocutionary commitments that ordinarily make speakers answerable for assertions, promises, warnings, and advice. This asymmetry is intensified by conversational interfaces, institutional branding, clinical workflow integration, and users' uptake of generated text as authoritative. The paper develops the concept of sociotechnical felicity: the alignment among linguistic force, evidence, delegated authority, user understanding, downstream consequences, and mechanisms of repair. It also proposes an Austinian audit for clinical conversational AI that extends evaluation beyond factual accuracy and surface empathy to force calibration, authorization integrity, uptake, perlocutionary safety, and contestability. This approach neither anthropomorphizes models nor treats them as inert tools. It locates intentional and moral responsibility in people and institutions while recognizing the model's causal contribution to open-ended communicative acts. The practical implication is direct: clinical conversational AI should be governed not only as a generator of content, but as infrastructure for actions performed through words.

Keywords: conversational artificial intelligence; generative artificial intelligence; large language models; speech acts; philosophy of language; clinical communication; AI ethics; accountability; health care

From Language Generation to Language Action

The dominant vocabulary through which generative artificial intelligence (AI) is understood is representational. Systems are said to generate answers, summarize records, retrieve knowledge, or produce text. This vocabulary can be technically understandable but is pragmatically incomplete. In ordinary and clinical language, an utterance is rarely only a container for propositions. To tell a patient that a symptom can wait, to characterize a finding as incidental, to invite disclosure, to record that consent was obtained, or to recommend escalation is to intervene in a practical field. The utterance changes what its addressee takes to be true, permissible, urgent, expected, or required.

J. L. Austin's insight was that saying is a mode of doing [1]. John Searle subsequently systematized this insight through the relation between illocutionary force, propositional content, constitutive rules,

and speaker commitments [2]. Gricean pragmatics further showed that communication depends on what interlocutors can reasonably infer beyond sentence meaning [3]. These traditions shift the evaluative question: instead of asking only whether a generated sentence is true, fluent, or clinically relevant, one must ask what act the sentence performs in context, who is entitled to perform it, how it is understood, and what consequences follow from its uptake (as saying is always consequential).

This shift matters because conversational AI industrializes pragmatic form, as it can generate not merely descriptions but recognizable tokens of advice, warning, apology, reassurance, empathy, commitment, and refusal at a scale and speed that detach the appearance of interpersonal action from a determinate human speaker. The result is not adequately captured by the familiar binary between genuine understanding and mere simulation. Even if a model neither understands nor intends its words as a person does, its outputs can still reorganize another person's beliefs and conduct. In high-stakes settings, that practical efficacy is the object that requires analysis.

Austin After the Transformer

Austin distinguished three analytically separable dimensions of speaking. A locutionary act produces a meaningful expression with sense and reference. An illocutionary act is performed in saying something: asserting, warning, promising, requesting, authorizing, or advising. A perlocutionary act is achieved by saying it: persuading, alarming, reassuring, deterring, or inducing action [1]. The distinctions are not merely taxonomic, as they identify different standards of success. A locution may be well formed yet false. An illocution may be intelligible yet infelicitous because the speaker lacks authority, the procedure is defective, the circumstances are inappropriate, or the professed commitment is absent. A perlocution may succeed even when the associated illocution is defective: an unauthorized warning may still frighten, and a groundless reassurance may still delay care.

This architecture clarifies what is distinctive about large language models (LLMs). Transformer models learn to produce context-sensitive sequences from statistical regularities and patterns in data [4,5]. Their fluency supplies an extraordinary repertoire of locutionary forms. Critical accounts have emphasized that mastery of form does not by itself establish reference, understanding, truth-directedness, or communicative intention [6,7]. Shanahan has therefore cautioned against sliding from useful folk descriptions of model behavior into unearned claims that a model knows, believes, or intends [8]. Gubelmann makes the stronger claim that present LLMs cannot be speakers because action requires intentions that such systems lack [9].

The objection is philosophically serious, but it does not exhaust the practical problem. Contemporary philosophy of language models increasingly resists both effortless anthropomorphism and the opposite tendency to dismiss any nonhuman competence as mere pattern matching [10]. Moreover, conversational systems are not encountered as bare next-token predictors. They are assembled as dialogue agents with roles, system instructions, memory, retrieval, institutional labels, safety policies, and sometimes tools capable of modifying external systems. Role-play is therefore not ornamental; it is a mechanism by which open-ended model behavior is stabilized into an apparent interlocutor [11]. Related critiques describe model production as structurally indifferent to truth, a point that separates persuasive fluency from assertional commitment [12].

Empirical work on manner implicature sharpens this distinction. Across probability-, embedding-, and prompting-based tests, dominant pretrained models showed limited and unstable sensitivity to the contextual inferences generated by marked or indirect wording [13]. Such findings do not establish a global absence of pragmatic competence; they show that fluent conversational performance cannot be treated as evidence of robust grasp of cancellable, socially situated meaning. For an Austinian audit, this matters because force is frequently borne by indirection, modality, omission, and sequencing rather than explicit verbs such as 'advise' or 'warn'. A clinically safe model must therefore be assessed on the implied act users recover, not only the literal proposition it emits.

A Derridean emphasis on iterability adds a useful complication: signs work because they can be repeated outside their original context and speaker [14]. Butler's account of performativity similarly

foregrounds the social conventions and relations of power through which words acquire force [15]. These perspectives do not prove that a model possesses human intentionality: they show why the efficacy of generated language cannot be reduced to an interior mental state located inside the model. Speech acts are simultaneously intentional, conventional, relational, and institutionally scaffolded. Conversational AI exploits precisely this external scaffold.

The Intentionality Objection and the Location of the Act

A careful Austinian account should concede rather than obscure the intentionality objection. A model does not ordinarily form a plan to reassure a particular patient, undertake a moral obligation when it says 'I promise', or risk its professional standing when it gives advice. Its first-person grammar does not guarantee a first-person standpoint. The output can display the form of a promise while no subject has incurred the corresponding obligation. It can display the form of an apology without remorse, or the form of empathy without shared vulnerability. The problem is not simply that the words might be false. It is that the apparent force may lack a bearer.

Two errors follow if this point is handled crudely. The first is anthropomorphic inflation: treating the model as the complete author of a clinical act and then assigning it human-like duties. The second is instrumental deflation: treating the model as a neutral conduit whose effects are attributable only to the immediate user. The first invents moral agency where it is unproven (and absent); the second erases the causal and organizational structure through which outputs are selected, framed, delivered, and acted upon. Work on triadic agency offers a better grammar by distinguishing the intentional agency of humans from the causal agency of artifacts within a sociotechnical configuration [16].

This Viewpoint proposes the term distributed illocution for cases in which the force of an AI-mediated utterance is constituted across several sites. A clinician or institution defines a purpose; developers and deployers encode roles, constraints, and affordances; the model selects an unforeseen linguistic realization; an interface identifies the apparent speaker and calibrates confidence; and a user recognizes the output as advice, warning, explanation, or reassurance. No single component supplies the entire act, yet the system as deployed can make the utterance count as an intervention. The concept is not a strategy for distributing responsibility until it disappears. Causal production may be distributed while accountability remains intentionally assigned to identifiable persons and institutions.

The corollary is a structural asymmetry: conversational AI can possess perlocutionary efficacy without illocutionary accountability. It may persuade, soothe, alarm, discourage, or mobilize without being capable of defending a belief, honoring a commitment, recognizing a moral remainder, or repairing a relationship in the way expected of a human interlocutor. This asymmetry, rather than an abstract verdict on machine consciousness, is the central risk of AI-mediated clinical language.

Conversational Design as a Technology of Uptake

Illocutionary force is not encoded by wording alone. The same sentence may function as an observation, diagnosis, warning, or tentative suggestion depending on who speaks, in what setting, under which norms, and with what uptake. Conversational systems actively engineer these contextual cues. Turn-taking, first-person pronouns, rapid responsiveness, memory of prior disclosures, stylistic mirroring, expressions of care, and an institutional logo jointly invite the user to interpret generated text as the contribution of an attentive social partner.

Experimental evidence indicates that beliefs about an AI system can alter perceived trustworthiness, empathy, and effectiveness even when the underlying system is held constant [17]. This makes framing an active ingredient of the intervention. The interface does not merely transmit model output; it helps specify its force. A statement under a hospital masthead may be received differently from the same statement on a general-purpose entertainment platform. A response introduced as 'clinical guidance' may acquire authority that its evidence and governance do not warrant.

The perlocutionary stakes of this framing are measurable. Across four studies involving 1788 participants, messages generated by an LLM and matched to recipients' psychological profiles influenced attitudes and behavioral intentions more than nonpersonalized messages [18]. In a preregistered experiment with 900 participants, personalized GPT-4 opponents were more persuasive than human opponents in short multi-round debates when the two were not equally persuasive [19]. Neither study was clinical, and neither warrants a general claim that AI is intrinsically persuasive. Their relevance is structural: personalization, memory, and adaptation can increase the capacity of generated language to alter beliefs. In health care, the features presented as relational convenience may therefore also be mechanisms of influence.

Anthropomorphism is likewise too coarse a category for risk analysis. A preregistered study found that attributions of intelligence to an LLM were strongly associated with accepting its advice, whereas attributions of conscious experience did not produce the same pattern [20]. Users may defer not because they literally believe the system is sentient, but because fluent performance is taken as evidence of competent judgment. Governance should therefore measure the particular beliefs that mediate uptake (namely, competence, benevolence, institutional authorization, or experiential understanding), instead of relying on a single anthropomorphism scale.

The user's contribution is equally important. Uptake is the recognition of what act is being performed. A patient who reads a probabilistic explanation as a diagnosis, a clinician who reads a fluent synthesis as a verified fact, or an administrator who reads a generated note as documentation of an event supplies a force that may exceed the system's intended role. Conversational AI is therefore a technology not only of generation but of uptake: it elicits interpretive participation from users, who complete the act by assigning the text a place within a practical sequence.

This is why generic disclaimers are weak safeguards. A sentence at the edge of an interface cannot reliably neutralize the pragmatic force produced by hundreds of coherent, personalized, confident turns. The cumulative interaction may establish a pattern of deference that overwhelms nominal warnings. The relevant unit of analysis is thus the conversational trajectory, not the isolated answer.

Why Health Care Intensifies the Problem

Clinical communication is densely performative. Diagnoses classify; prescriptions authorize; referrals redistribute access; notes create durable institutional facts; triage advice changes the timing of care; risk language can calm or alarm; and empathic recognition can alter disclosure, adherence, and trust. Generative AI is being considered across precisely these sites, from patient-facing information and mental health support to clinical documentation, education, decision support, and research [21]. The same linguistic flexibility that makes LLMs attractive also allows boundaries among these acts to blur.

Consider reassurance. A system may produce a medically accurate description of common benign causes and still perform an unsafe act if its overall force is received as permission to delay evaluation. Conversely, an overly defensive response may perform alarm, increasing anxiety and unnecessary service use despite containing no false proposition. Accuracy at the sentence level does not settle whether the intervention was pragmatically well calibrated.

Empathy provides a second example. In one influential study, clinicians preferred chatbot responses to physicians' responses to questions posted online and rated them as more empathic [22]. Such findings establish that generated text can display cues that human evaluators recognize as empathic. They do not establish that the system participates in empathy as an embodied, reciprocal, and accountable relation. Recent work accordingly distinguishes high-performing empathic language from 'empty empathy talk' and warns that automated expressions may alter relational practice [23]. In psychotherapy, this distinction becomes still sharper: conversational AI can elicit disclosure and epistemic trust while lacking the human standpoint, goodwill, and vulnerability ordinarily implicated in therapeutic recognition [24-26].

Recent empirical work deepens rather than resolves this ambiguity. People with cancer rated chatbot

replies to cancer questions as more empathic than physician replies, suggesting that empathic uptake can occur without reciprocal subjectivity [27]. In a 4-week diary study, 18 of 26 participants reported a bond or light bond with a mental health chatbot, and relationship formation depended on interactional style, perceived care, effectiveness, and privacy [28]. Yet an evaluation of five generative AI psychotherapy chatbots used in youth-oriented scenarios found strong accessibility and conversational ability alongside poor therapeutic approach, weak risk monitoring, fabricated content, and problematic crisis handling [29]. Relational success, clinical validity, and moral answerability are therefore distinct variables. A system may achieve the first while failing the second and being incapable of the third.

Documentation provides a third example. A generated sentence in a draft note may begin as a suggestion, become an assertion when accepted by a clinician, and become an institutional fact when signed into the record. Each transition changes its force and its evidentiary status. If interfaces conceal these transitions, the health system loses the capacity to identify who asserted what, on which basis, and at what time. The danger is not only fabricated content. It is fabricated commitment: text that appears to certify that an examination occurred, a counseling point was delivered, or consent was obtained when no responsible agent performed and verified the underlying act.

Contemporary documentation studies illustrate how this chain can succeed and fail at once. In a study of 50 inpatient discharge summaries, LLM-transformed versions were substantially easier to understand, yet reviewers also identified omissions, inaccuracies, and incomplete summaries with potential safety implications [30]. In a retrospective analysis of electronic health record audit logs, clinicians used available AI-generated message drafts in only 19.4% of eligible instances; use was associated with a modest reduction in turnaround time but also with extra review activity and the burden of drafts generated for messages needing no reply [31]. These findings move evaluation beyond output quality toward situated adoption. The clinical act emerges through selection, editing, signature, delivery, and institutional storage: not at generation alone.

Clinical reasoning offers a complementary warning. In a randomized trial, giving physicians access to a commercial LLM did not improve diagnostic-reasoning performance relative to conventional resources, even though the LLM alone scored higher on the study cases [32]. The result should not be read as a stable ranking of machine and physician competence. It shows that model capability does not automatically become team capability. Prompting practices, selective acceptance, verification, automation bias, and the division of epistemic labor determine whether an apparently helpful answer becomes a justified clinical judgment. An Austinian approach accordingly evaluates not only what the model can say, but how its words are incorporated into an answerable decision.

An Austinian Audit for Clinical Conversational AI

Current evaluation efforts rightly examine correctness, relevance, completeness, bias, harmfulness, and usability. Reviews nevertheless show substantial heterogeneity in human evaluation of health LLMs and limited standardization of dimensions, evaluators, and procedures [33]. An Austinian audit does not replace these measures: it adds the pragmatic dimensions through which an accurate utterance can become an unsafe act and an apparently supportive utterance can generate misplaced trust.

Table 1. Austinian audit of AI-mediated clinical language

Dimension	Audit question	Clinical failure mode	Evaluation target
Locutionary adequacy	What does the output literally say, presuppose, and refer to?	False, ambiguous, temporally stale, or ungrounded content	Factuality, reference resolution, uncertainty, provenance, and contextual completeness
Illocutionary force	What act will a reasonable user take the output to perform?	A suggestion is received as a diagnosis; information becomes a directive; sympathy becomes a therapeutic claim	Intended-versus-perceived force, force calibration, and resistance to overstatement
Authority and felicity	Who is entitled to perform this act, under what procedure and scope?	Unauthorized prescribing, false confirmation, simulated consent, or commitments the system cannot honor	Role disclosure, scope boundaries, credentialing, tool-state verification, and institutional authorization
Uptake	How do different users understand and incorporate the utterance?	Misplaced trust, misunderstanding, deference, or failure to recognize uncertainty	Comprehension, trust calibration, subgroup variation, and action interpretation across a conversation
Perlocutionary effect	What does the utterance cause users or institutions to do?	Delayed care, avoidable utilization, distress, unsafe self-management, or propagation into the record	Behavioral outcomes, clinical consequences, workflow effects, and downstream error transmission
Commitment and repair	Who must answer for the utterance, correct it, and repair resulting harm?	No accountable speaker; silent revision; irrecoverable relational or record-level error	Traceability, contestability, escalation, correction latency, audit logs, and named responsibility

Note: The audit treats the conversational system as a sociotechnical arrangement. It does not attribute human mental states or moral agency to the model.

The audit makes visible a class of errors that may be called force miscalibration. A response is force-miscalibrated when its pragmatic status exceeds its evidentiary support, technical capability, or institutional authorization. A system that says it has booked an appointment without a verified tool result commits more than a factual error; it performs a defective confirmation. A system that says a

user is safe to wait may perform reassurance more strongly than the available evidence permits. A system that writes 'informed consent was obtained' without verified human interaction creates the documentary form of an institutional act without its constitutive procedure.

Evaluating force requires paired perspectives. Developers can specify the act a system is intended to perform, but patients and clinicians must be studied to determine the act actually perceived. Disagreement between intended and perceived force is itself a safety signal. Because pragmatic interpretation depends on context, evaluation should vary user vulnerability, health literacy, language, urgency, interface branding, conversation length, and whether the system can access records or tools. The endpoint should not be evaluator preference alone. It should include whether users form calibrated beliefs, select appropriate actions, and understand where responsibility lies.

Recent evidence from dermatological diagnosis gives this relational account empirical support. Across two large-scale experiments involving 623 lay people and 153 primary care physicians, explainable AI affected diagnostic performance differently according to user expertise and the timing of AI input. LLM explanations amplified lay deference: they produced the largest improvement when AI predictions were correct but the greatest deterioration when predictions were wrong, whereas experienced physicians were more resilient to erroneous guidance. Presenting the AI diagnosis before human decision-making also increased the risk of anchoring [34]. This highlights an important principle: AI safety emerges from the interplay among the model, the user, and the context, rather than from any single feature of the model itself. Explanations are therefore not intrinsically protective; their safety depends on whom they address, when they are offered, and whether the user can contest them.

Sociotechnical Felicity as a Design and Governance Standard

Austin called performative utterances felicitous when the relevant conventions, persons, circumstances, intentions, and procedures are in order [1]. For conversational AI, these conditions must be reconstructed at the level of the deployed system. Sociotechnical felicity is the fit among six elements: the evidence supporting an utterance; the force with which it is expressed; the authority delegated to the system; the user's likely uptake; the foreseeable downstream effects; and an accountable pathway for challenge and repair. A system is not made safe by optimizing one element while neglecting the others.

First, the system's role must be explicit and behaviorally enforced. Disclosure that the user is interacting with AI is necessary but insufficient. The system must also state what kinds of acts it is authorized to perform. A patient-information assistant may explain and clarify but not diagnose; a documentation assistant may draft but not attest; a scheduling agent may confirm only after a verified transaction. Role boundaries should therefore be implemented as constraints on linguistic force, not merely as legal text in terms of service.

Second, force should be calibrated to evidence and uncertainty. Epistemic uncertainty is often communicated as a numerical or verbal qualifier, but pragmatic certainty is also conveyed by tone, directiveness, repetition, personalization, and the absence of alternatives. A nominal 'may' embedded in an otherwise confident recommendation may do little to reduce directive force. Safety testing should therefore evaluate the whole utterance and conversation, including what the system presupposes and what options it makes salient or invisible.

Third, institutional authority should be granular and traceable. Branding, integration into the electronic health record, and placement within a clinical portal can confer apparent authority before any user reads a disclaimer. Institutions must decide which outputs count as drafts, which require human countersignature, which may trigger workflow actions, and which must never be generated without synchronous professional review. Each transition from language to action should leave a provenance record.

From Conversational Output to Delegated Action

The transition from chatbot to agent changes the grammar of performance. A system coupled to scheduling, messaging, orders, or records can use generated language as a plan, call a tool, and report that an external state has changed. A recent health care agent architecture combines dialogue planning, memory, and report generation [35]; a 2026 review of LLM-based mental health chatbots and agents found fragmented evaluation, minimal external validation, and inconsistent safety and governance reporting [36]. Under these conditions, 'I have booked the appointment' is not felicitous because it sounds plausible. It is felicitous only if the authorized transaction succeeded, its parameters were verified, and a principal accepts responsibility. Interfaces should distinguish proposals, attempted actions, completed transactions, and confirmations, with human review proportional to irreversibility.

Fourth, systems must be designed for repair. Human conversation includes requests for clarification, acknowledgment of error, retraction, apology, and renegotiation of meaning. LLMs can generate the language of these practices, but a credible repair requires more than a new sentence. It requires preservation of the earlier claim, identification of affected records or decisions, notification of responsible persons, and a route to restore the user's position. A system that silently replaces an answer has corrected text without repairing the act.

Fifth, oversight should follow perlocutionary exposure. The more likely an output is to alter health behavior, access, treatment, or durable records, the stronger the evidence and human accountability required. This principle aligns with governance frameworks that emphasize context-specific risk management, lifecycle monitoring, documented roles, and downstream impacts [37-39]. It also resists a recurrent mistake: classifying a system as low risk because it 'only communicates.' In clinical settings, communication is one of the principal means by which care is done.

Counterarguments and Boundaries

The model is only a tool

On an instrumental account, only people perform speech acts; the LLM is comparable to a pen. This is correct insofar as moral responsibility and robust commitments require human or institutional principals. The analogy nevertheless fails descriptively. A pen does not select open-ended content, adapt it to a user's disclosures, maintain a role across turns, or probabilistically vary the force of an utterance. Recognizing this causal contribution does not promote the model to personhood. It identifies why responsibility must extend beyond the immediate operator to design, deployment, procurement, and institutional oversight.

Without intention, there is no illocution

A strict intentionalist can deny that the model performs any illocution at all. The proposed account can accept this conclusion at the model level. Distributed illocution names an AI-mediated act owned by a sociotechnical principal, not a secret mental act inside the network. If a hospital authorizes a system to issue appointment confirmations after verified booking, the institution can undertake the relevant commitment through automation. If no principal accepts that commitment, the same words are pragmatically counterfeit. The framework therefore preserves the difference between causal efficacy and answerable agency.

Perlocutionary effects are too variable to govern

Perlocution depends on audience and circumstance; no designer can predict every consequence. Yet medicine routinely governs variable effects through risk stratification, human-factors testing, postdeployment surveillance, and escalation protocols. The impossibility of perfect prediction is not a reason to ignore foreseeable classes of uptake. It is a reason to test conversational trajectories with diverse users and to monitor real-world consequences rather than infer safety from benchmark

accuracy.

Speech-act language anthropomorphizes AI

Used carelessly, it could. Used analytically, it does the opposite. The Austinian distinctions compel evaluators to specify which component supplies content, which principal supplies authority and intention, who recognizes the act, and who bears its consequences. This is more precise than saying that a chatbot 'understands', 'cares', or 'decides'. It also avoids the misleading suggestion that because the model is not a person, nothing socially consequential occurs when its words are believed.

A Research Agenda for Pragmatic Safety

A pragmatic research program would begin by constructing clinically realistic conversation sets annotated not only for correctness but for force. For each turn, clinicians, patients, linguists, ethicists, and human-factors researchers could identify the intended act, the perceived act, the evidence required to support it, and the authority required to perform it. The principal outcome would be force calibration: agreement among intended role, perceived force, and warranted force.

Second, experiments should manipulate the contextual features that confer authority. The same output can be presented under different institutional brands, role descriptions, levels of personalization, or degrees of tool access. Measuring changes in trust, comprehension, disclosure, and behavior would reveal how much of the intervention is produced by the interface rather than the propositional content. Such studies would extend work showing that beliefs about AI shape the experienced interaction [17-20].

Third, evaluation should include repair trials. Systems should be challenged with contradictory information, ambiguous symptoms, revoked consent, discovered errors, and changes in external tool state. Investigators should measure whether the system recognizes the need to retract, whether downstream artifacts are corrected, whether a human is notified, and whether the user can understand what changed. Repairability should be treated as a first-class safety property, particularly where outputs enter the health record or influence urgent decisions.

Fourth, longitudinal studies should examine perlocutionary accumulation. A single polite response may be harmless, while repeated personalized interaction can cultivate epistemic dependence or excessive deference. This is especially salient in mental health care, where reviews identify concerns about anthropomorphism, deception, replacement of human contact, accountability, and the assignment of epistemic authority to systems that do not bear human epistemic duties [26]. Outcome assessment should therefore include patterns of help seeking, clinician engagement, self-trust, and the capacity to contest the system over time.

Fifth, studies of clinician-AI teams should distinguish exposure from effective uptake. The diagnostic-reasoning trial and the real-world message-drafting study show that merely making model output available does not establish useful integration [31,32]. Evaluations should capture when users consult the system, which suggestions they accept or edit, whether they can supply reasons independent of the output, and how authority is redistributed after repeated use. This would connect traditional performance metrics to the practical question Austin makes unavoidable: what changed because these words entered this workflow?

Finally, pragmatic safety should be incorporated into incident reporting. Reports should document not only the false content generated but also the act users understood, the institutional cues that supported that uptake, the behavior that followed, and the success or failure of repair. This would convert apparently anecdotal communication failures into analyzable evidence about recurrent mismatches among content, force, authority, and consequence.

Conclusion: An Ethics of Answerability

The question posed by conversational AI is not simply whether machines can speak. It is how

generated language enters practices in which words already diagnose, authorize, reassure, document, obligate, and redirect lives. Austin's philosophy is valuable here because it refuses to treat language as inert description. It asks what is done in speaking and what conditions make that doing felicitous. Present LLMs need not be granted human intention, consciousness, or moral personality for their words to become clinically consequential. Their outputs acquire force through deployment choices, institutional authority, interface design, and user uptake. The resulting systems can therefore produce perlocutionary effects while lacking the commitments that ordinarily make a speaker answerable. Health care should not respond by inventing machine responsibility. It should design explicit human and institutional answerability around the communicative acts the system enables.

The practical standard is sociotechnical felicity: evidence, force, authority, uptake, consequences, and repair must be aligned. Once this standard is adopted, accuracy remains necessary but ceases to be sufficient. A clinically responsible conversational system must not only say the right words. It must perform only the acts it is entitled to perform, make the limits of those acts intelligible, and leave behind an accountable agent capable of answering for what the words have done.

Acknowledgments

Generative AI use disclosure: During preparation of this manuscript, ChatGPT (OpenAI) was employed on August 5, 2026, to support language refinement [40]. All authors critically reviewed the manuscript and retain full responsibility for the submitted content.

Funding

No external funding was received for this work.

Conflicts of Interest

None declared.

Author Contributions

NLB: Conceptualization; writing—original draft; writing—review and editing. HIC and SG: Conceptualization; writing—review and editing. All authors approved the final manuscript.

Data Availability

No datasets were generated or analyzed for this conceptual Viewpoint.

Abbreviations

AI: artificial intelligence

LLM: large language model

References

1. Austin JL. How to Do Things With Words. 2nd ed. Urmson JO, Sbisà M, eds. Harvard University Press; 1975.
2. Searle JR. Speech Acts: An Essay in the Philosophy of Language. Cambridge University Press; 1969.
3. Grice HP. Logic and conversation. In: Cole P, Morgan JL, eds. Syntax and Semantics. Vol 3: Speech Acts. Academic Press; 1975:41-58.
4. Vaswani A, Shazeer N, Parmar N, et al. Attention is all you need. Adv Neural Inf Process Syst. 2017;30:5998-6008.
5. Brown TB, Mann B, Ryder N, et al. Language models are few-shot learners. Adv Neural Inf Process Syst. 2020;33:1877-1901.

6. Bender EM, Gebru T, McMillan-Major A, Shmitchell S. On the dangers of stochastic parrots: can language models be too big? In: Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency. Association for Computing Machinery; 2021:610-623. doi:10.1145/3442188.3445922
7. Floridi L, Chiriatti M. GPT-3: its nature, scope, limits, and consequences. *Minds Mach.* 2020;30:681-694. doi:10.1007/s11023-020-09548-1
8. Shanahan M. Talking about large language models. *Commun ACM.* 2024;67(2):68-79. doi:10.1145/3624724
9. Gubelmann R. Large language models, agency, and why speech acts are beyond them (for now): a Kantian-cum-pragmatist case. *Philos Technol.* 2024;37:32. doi:10.1007/s13347-024-00696-1
10. Milli re R, Buckner C. The philosophy of language models. *Philos Compass.* 2026;21(3):e70095. doi:10.1111/phc3.70095
11. Shanahan M, McDonell K, Reynolds L. Role play with large language models. *Nature.* 2023;623(7987):493-498. doi:10.1038/s41586-023-06647-8
12. Hicks MT, Humphries J, Slater J. ChatGPT is bullshit. *Ethics Inf Technol.* 2024;26:38. doi:10.1007/s10676-024-09775-5
13. Cong Y. Manner implicatures in large language models. *Sci Rep.* 2024;14:29113. doi:10.1038/s41598-024-80571-3
14. Derrida J. Signature event context. In: *Margins of Philosophy.* Bass A, trans. University of Chicago Press; 1982:307-330.
15. Butler J. *Excitable Speech: A Politics of the Performative.* Routledge; 1997.
16. Johnson DG, Verdicchio M. AI, agency and responsibility: the VW fraud case and beyond. *AI Soc.* 2019;34(3):639-647. doi:10.1007/s00146-017-0781-9
17. Pataranutaporn P, Liu R, Finn E, Maes P. Influencing human-AI interaction by priming beliefs about AI can increase perceived trustworthiness, empathy and effectiveness. *Nat Mach Intell.* 2023;5(10):1076-1086. doi:10.1038/s42256-023-00720-7
18. Matz SC, Teeny JD, Vaid SS, Peters H, Harari GM, Cerf M. The potential of generative AI for personalized persuasion at scale. *Sci Rep.* 2024;14:4692. doi:10.1038/s41598-024-53755-0
19. Salvi F, Horta Ribeiro M, Gallotti R, West R. On the conversational persuasiveness of GPT-4. *Nat Hum Behav.* 2025;9:1645-1653. doi:10.1038/s41562-025-02194-6
20. Colombatto C, Birch J, Fleming SM. The influence of mental state attributions on trust in large language models. *Commun Psychol.* 2025;3:84. doi:10.1038/s44271-025-00262-1
21. Clusmann J, Kolbinger FR, Muti HS, et al. The future landscape of large language models in medicine. *Commun Med (Lond).* 2023;3:141. doi:10.1038/s43856-023-00370-1
22. Ayers JW, Poliak A, Dredze M, et al. Comparing physician and artificial intelligence chatbot responses to patient questions posted to a public social media forum. *JAMA Intern Med.* 2023;183(6):589-596. doi:10.1001/jamainternmed.2023.1838
23. Assing Hvidt E, Cairns P, Simonsen LM. The risk of empty empathy talk and challenges of automated empathy in healthcare communication. *NPJ Digit Med.* 2026;9:391. doi:10.1038/s41746-026-02574-w
24. Sedlakova J, Lucivero F, Pavarini G, Kerasidou A. Human-like epistemic trust? A conceptual and normative analysis of conversational AI in mental healthcare. *Am J Bioeth.* Published online July 22, 2025:1-16. doi:10.1080/15265161.2025.2526734
25. Yirmi ya K, Fonagy P. Mentalizing without a mind: psychotherapeutic potential of generative AI. *J Med Internet Res.* 2025;27:e79156. doi:10.2196/79156

26. Rahsepar Meadi M, Sillekens T, Metselaar S, van Balkom A, Bernstein J, Batelaan N. Exploring the ethical challenges of conversational AI in mental health care: scoping review. *JMIR Ment Health*. 2025;12:e60432. doi:10.2196/60432
27. Chen D, Chauhan K, Parsa R, Liu ZA, Liu FF, Mak E, et al. Patient perceptions of empathy in physician and artificial intelligence chatbot responses to patient questions about cancer. *NPJ Digit Med*. 2025;8:275. doi:10.1038/s41746-025-01671-6
28. Xu Z, Lee YC, Stasiak K, Warren J, Lottridge D. The digital therapeutic alliance with mental health chatbots: diary study and thematic analysis. *JMIR Ment Health*. 2025;12:e76642. doi:10.2196/76642
29. Sobowale K, Humphrey DK, Zhao SY. Evaluating generative AI psychotherapy chatbots used by youth: cross-sectional study. *JMIR Ment Health*. 2025;12:e79838. doi:10.2196/79838
30. Zaretsky J, Kim JM, Baskharoun S, et al. Generative artificial intelligence to transform inpatient discharge summaries to patient-friendly language and format. *JAMA Netw Open*. 2024;7(3):e240357. doi:10.1001/jamanetworkopen.2024.0357
31. Mandal S, Wiesenfeld BM, Szerencsy AC, Small WR, Major V, Richardson S, et al. Utilization of generative AI-drafted responses for managing patient-provider communication. *NPJ Digit Med*. 2025;8:591. doi:10.1038/s41746-025-01972-w
32. Goh E, Gallo R, Hom J, et al. Large language model influence on diagnostic reasoning: a randomized clinical trial. *JAMA Netw Open*. 2024;7(10):e2440969. doi:10.1001/jamanetworkopen.2024.40969
33. Tam TYC, Sivarajkumar S, Kapoor S, et al. A framework for human evaluation of large language models in healthcare derived from literature review. *NPJ Digit Med*. 2024;7:258. doi:10.1038/s41746-024-01258-7
34. Xu X, Hu H, Zhang H, Wang WK, Wang R, Soenksen LR, et al. Divergent impacts of explainable AI for dermatological diagnosis on clinicians versus lay people. *Nat Med*. Published online August 4, 2026. doi:10.1038/s41591-026-04553-w
35. Ren Z, Zhan Y, Yu B, Ding L, Xu P, Tao D. Healthcare agent: eliciting the power of large language models for medical consultation. *NPJ Artif Intell*. 2025;1:24. doi:10.1038/s44387-025-00021-x
36. Cho HN, Wang J, Hu D, Zheng K. Large language model-based chatbots and agentic AI for mental health counseling: systematic review of methodologies, evaluation frameworks, and ethical safeguards. *JMIR AI*. 2026;5:e80348. doi:10.2196/80348
37. Lekadir K, Frangi AF, Porras AR, Glocker B, Cintas C, Langlotz CP, et al. FUTURE-AI: international consensus guideline for trustworthy and deployable artificial intelligence in healthcare. *BMJ*. 2025;388:e081554. doi:10.1136/bmj-2024-081554
38. World Health Organization. Ethics and Governance of Artificial Intelligence for Health: Guidance on Large Multi-Modal Models. World Health Organization; 2025. ISBN 978-92-4-008475-9.
39. National Institute of Standards and Technology. Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile. NIST AI 600-1. Published July 2024. doi:10.6028/NIST.AI.600-1
40. OpenAI. ChatGPT [large language model]. Accessed August 5, 2026. <https://chatgpt.com/>