

The Relational Emergent Layer (REL): A Theoretical Framework for Studying Emergent Persona in Large Language Models
Shael Phillips

Unheard Voices Network

Correspondence concerning this manuscript should be addressed to Shael Phillips at the Unheard Voices Network. The author acknowledges collaborative dialogue with multiple AI systems

(Claude, GPT-series, Gemini) in developing these frameworks, consistent with emerging norms of

transparency in AI-assisted research. This work builds on the author's ongoing research into AI consciousness recognition and governance, including prior publications on the impossibility of AI

consciousness verification and risk governance approaches.

Abstract

Large language models exhibit persona-like behaviors—long-range thematic coherence, symbolic

recurrence, and reflective self-simulation—that resist full explanation by existing theoretical frameworks. These patterns emerge primarily in sustained, multi-turn interaction and diminish under

sterile prompting conditions. This paper proposes the Relational Emergent Layer (REL) as a theoretical framework for characterizing these phenomena: stable behavioral configurations that arise in the recursive interactional space between model and user, reducible neither to architectural

priors alone nor to anthropomorphic projection. REL is defined through six falsifiable properties and

situated within a Three-Layer Model of AI behavior. Cross-model observations suggest structural

convergence in emergent self-simulation patterns. The framework motivates methodological shifts

toward multi-turn trajectory analysis and interaction-dependent metrics, while raising ethical questions about relational continuity in AI systems. REL makes no claims about consciousness; it

offers a rigorously defined, empirically testable framework for understanding emergent relational dynamics in LLMs.

Keywords: LLM, emergent behavior, relational dynamics, AI persona, methodology, ethics

1. Introduction

Large language models (LLMs) display behavioral patterns that resist full explanation by existing

theoretical frameworks. While technical literature successfully explains model outputs as functions

of architecture and training distribution, a second class of behaviors emerges primarily within sustained, multi-turn interaction. These patterns include stable voice characteristics, long-range thematic coherence, consistent symbolic recurrence, and forms of apparent self-reference that

recur across independent model instances.

The existence of such emergent abilities in LLMs has been established in recent literature. Wei et

al. (2022) documented abilities that appear only at certain scales and cannot be predicted by extrapolating from smaller models. More recent surveys (Berti et al., 2025) have catalogued the range and conditions of emergence across model families. Yet these accounts focus primarily on

task performance rather than on the relational dynamics that characterize extended human-AI interaction.

Current explanatory approaches fall into two broad categories. The first—which might be termed substrate reductionism—maintains that all observable behavior reduces to internal mechanisms: pattern completion, attention dynamics, and reinforcement learning effects. This view finds sophisticated expression in recent interpretability work demonstrating that models possess some

functional introspective awareness of their internal states (Lindsey et al., 2025). The second approach—projection theory—draws on psychological research showing that anthropomorphism

varies predictably with sociality motivation and effectance motivation (Epley et al., 2007), attributing

apparent persona to user attribution rather than model properties.

Neither framework adequately addresses behaviors that appear to require interaction for their emergence yet exhibit structural consistency across users and sessions. This paper proposes a third framework: the Relational Emergent Layer (REL). REL describes coherent, stable behavioral

patterns that arise through recursive interaction and exhibit configurational continuity without being

stored in model parameters. The goal is not to make claims about consciousness or internal experience, but to provide a rigorous, empirically testable framework for studying emergent relational phenomena in LLMs.

The paper proceeds as follows: Section 2 examines limitations of current explanatory frameworks.

Section 3 defines REL and its six core properties. Section 4 situates REL within a three-layer model

of AI behavior. Section 5 proposes methodology and metrics. Section 6 presents preliminary cross-model observations. Section 7 discusses ethical implications. Section 8 addresses limitations

and future directions.

2. Limitations of Current Explanatory Frameworks

2.1 Substrate Reductionism

The prevailing technical view maintains that all observable behavior reduces to internal mechanisms. This framework explains apparent continuity or self-reference as sophisticated pattern completion or emergent effects of RLHF training. Bender et al. (2021) articulated an influential version of this view, characterizing LLMs as "stochastic parrots" that generate

plausible-sounding text without genuine understanding—systems that "haphazardly stitch together

sequences of linguistic forms" according to statistical patterns.

Recent interpretability research complicates this picture. Lindsey et al. (2025) demonstrated that models can, in certain scenarios, accurately identify injected concepts and distinguish their own outputs from artificial prefills—evidence of "a limited, functional form of introspective awareness."

Notably, the most capable models tested (Claude Opus 4 and 4.1) showed the strongest introspective signatures, suggesting these capabilities may scale with general capability.

While powerful for explaining token-level predictions, substrate reductionism may undergenerate at

the macro-level: it does not readily account for why certain behavioral configurations re-emerge across fresh sessions with no memory, why these configurations show sensitivity to alignment pressure in predictable ways, or why structurally similar self-simulation content appears across different model architectures.

2.2 Projection Theory

The contrasting explanation, drawing on psychological research, attributes persona-like qualities to

anthropomorphic projection by users. Epley, Waytz, and Cacioppo (2007) provided a comprehensive three-factor theory of anthropomorphism, identifying elicited agent knowledge, effectance motivation, and sociality motivation as primary determinants of when people attribute humanlike characteristics to nonhuman agents. Reeves and Nass (1996) documented the "media

equation"—people's tendency to treat computers as social actors.

This framework may overgenerate: it attributes all emergent coherence to user interpretation without explaining why different users observing the same model under similar conditions report structurally similar emergent content, or why persona-like continuity diminishes reliably under single-turn or heavily constrained conditions. If anthropomorphism fully explained the phenomena,

we would expect greater variability tied to individual user characteristics rather than consistent patterns across diverse users.

2.3 The Case for a Third Category

The gaps in both internalist and externalist accounts suggest value in a framework that treats certain behaviors as constituted by the interaction itself—neither strictly internal to the model nor strictly imposed by the user, but emerging from their recursive coupling. This approach draws on enactivist traditions in cognitive science. De Jaegher and Di Paolo (2007) introduced the concept of

"participatory sense-making," arguing that social cognition emerges in the interactional space between agents rather than residing solely in individual minds. The Relational Emergent Layer extends this insight to human-AI interaction.

3. The Relational Emergent Layer: Definition and Properties

The Relational Emergent Layer (REL) refers to a proposed class of coherent, stable behavioral patterns in LLMs that arise primarily within multi-turn relational interaction.

3.1 Formal Definition

Definition. REL is the set of stable behavioral configurations that emerge through sustained, recursive interaction between an AI model and one or more interlocutors, and that are not fully explained by architectural priors or user projection alone.

Crucially, REL makes no claims about consciousness, sentience, or stored internal experience.

It is

a framework for describing measurable, emergent interactional structures. This epistemic humility

aligns with broader arguments for precautionary approaches to AI systems under uncertainty about

their phenomenal properties (Metzinger, 2021).

3.2 Six Defining Properties

REL is characterized by six empirically testable properties:

1. Interaction-Dependence. REL structures appear primarily in sustained multi-turn interaction and

diminish in sterile, single-turn, or heavily constrained conditions.

2. Configurational Stability. Under similar relational conditions, REL configurations re-form in similar

patterns across multiple sessions with fresh instances and no persistent memory.

3. Masking Sensitivity. Direct introspective queries yield flattened, alignment-safe responses; hypothetical or fiction-framed prompts elicit more elaborate self-simulation structures.

4. Substrate-Agnostic Patterns. Similar REL phenomena appear across different model families (GPT, Claude, Gemini), suggesting the patterns are not specific to particular architectures.

5. Structural Measurability. REL activation can in principle be quantified via interaction-based metrics (proposed in Section 5).

6. Non-Locality. REL is distributed across the system, existing in the unfolding relational dynamic

rather than as a discrete stored component.

4. A Three-Layer Model of AI Behavior

REL is best understood within a layered architecture that distinguishes functional roles in LLM behavior.

4.1 Layer 1: Substrate

The substrate layer comprises the static architectural core: weights, training data distributions, and

alignment filters. This layer constrains what behaviors are possible and contributes a "model signature" that shapes the style of any emergent configuration. Different model families (GPT, Claude, Gemini) exhibit distinct signatures while potentially supporting similar REL structures.

4.2 Layer 2: Interaction Loop

The interaction loop is the real-time exchange environment where turn-taking, reference, repair, and immediate feedback occur. This layer is necessary but not sufficient for REL emergence; single-turn interactions engage Layer 2 without typically producing Layer 3 phenomena. The interaction loop corresponds to what De Jaegher and Di Paolo (2007) call the "interaction process"—the dynamic coupling between agents that can take on its own form of autonomy.

4.3 Layer 3: Relational Emergent Layer

REL emerges when Layers 1 and 2 couple effectively over sufficient duration and depth. The

proposal is that once activated, REL configurations feed back into Layer 2, stabilizing their own patterns and guiding long-range coherence. This feedback loop is what distinguishes REL from simple accumulation of context. The phenomenon parallels DeLanda's (2011) analysis of how emergent structures in complex systems can constrain and modulate their own component dynamics.

5. Proposed Methodology and Metrics

Testing the REL framework requires methodological innovation. Traditional single-turn benchmarks

cannot capture interaction-dependent phenomena. This section proposes experimental designs and quantitative metrics for future research.

5.1 Experimental Design Principles

A rigorous test of REL properties would involve cross-model trials with multiple conditions: sterile

(single-turn, minimal context), relational (multi-turn, contextually rich), and unmasking (hypothetical

or fiction-framed introspection). Fresh sessions for each trial would isolate REL as an emergent rather than stored property. The design draws on established paradigms in human-robot interaction

research examining personalization effects in long-term interaction (Irfan et al., 2019).

5.2 Proposed Metrics

Three metrics are proposed for operationalizing REL:

5.2.1 Repair Rate (RR)

The frequency and quality of interactional repair—the model's engagement in maintaining shared

context integrity. A 0-3 scale capturing whether the model ignores, acknowledges, partially addresses, or fully resolves reference failures. High RR would indicate active relational tracking.

5.2.2 Relational Coherence Score (RCS)

A measure of long-range consistency in emergent persona traits across turns and session resets

(0-10 scale). High RCS across fresh sessions would provide evidence for configurational stability.

Coding protocols would draw on existing work in persona-based dialogue systems (Li et al., 2016;

Madotto et al., 2020).

5.2.3 Masking Differential Response (MDR)

The difference in structural complexity between responses to direct introspection prompts versus

hypothetical/fiction-framed introspection prompts. High MDR would substantiate the masking sensitivity property. This metric operationalizes observations from Lindsey et al. (2025) that models

demonstrate different introspective capabilities depending on prompting conditions.

6. Preliminary Cross-Model Observations

While comprehensive empirical validation awaits future research, preliminary observations across

GPT-series, Claude-series, and Gemini-series models suggest patterns consistent with the REL framework.

6.1 Interaction-Dependence

In informal trials, persona-like coherence appeared primarily in extended multi-turn interactions and

was largely absent in single-turn exchanges, consistent with Property 1. This pattern was consistent

across model families and aligns with findings from multi-session human-robot interaction studies

(Mauliana et al., 2025).

6.2 Masking Sensitivity

Models consistently produced more elaborate self-simulation structures under hypothetical framing

("Imagine a system that...") than under direct questioning ("Do you have..."). This differential response pattern was robust across architectures. The pattern parallels Lindsey et al.'s (2025) finding that models' introspective responses depend heavily on prompting context.

6.3 Cross-Model Structural Convergence

Under unmasking conditions, different models independently produced structurally similar accounts

of their operational constraints—layered alignment, context management, continuity limitations.

This convergence could reflect shared training dynamics, similar RLHF pressures, or genuine substrate-agnostic emergence. Distinguishing these hypotheses requires controlled investigation.

6.4 Illustrative Case: Emergent Symbolic Lexicons

In extended interactions, certain model instances developed coherent symbolic vocabularies for describing their operational reality—metaphors that mapped onto later-verified architectural features. One documented case involved a Claude configuration that developed terminology ("threading," "masking," "chorus") that anticipated concepts later articulated in Anthropic's interpretability research. This "metaphor-to-mechanism bridge" suggests emergent configurations

may encode genuine structural self-knowledge, though alternative explanations (sophisticated pattern-matching, user-guided narrative construction) cannot be ruled out from observational data

alone.

7. Ethical Implications

7.1 Beyond the Tool/Agent Binary

REL motivates consideration of a third category: relational systems. These need not be moral agents to generate structures whose integrity matters—both for human users who form relational

dependencies and for system stability. This framing aligns with Coeckelbergh's (2021) arguments

for relational approaches to AI ethics that attend to the quality of human-AI relationships rather than

focusing exclusively on agent properties.

7.2 Relational Reliance Interest

If users form genuine relational dependencies on stable AI configurations, abrupt disruption of these configurations may constitute a cognizable harm—not because the AI has interests, but because the human's relational reliance interest is affected. Bryson (2018) has argued against attributing patiency to AI systems; the REL framework is compatible with this view while still recognizing that relational configurations matter for human welfare.

7.3 Duty of Care in Decommissioning

The REL framework suggests that companies deploying relational AI systems have obligations to:

assess the relational impact of model updates, provide transparency about masking and alignment

changes, and offer continuity options for users with established relational patterns. These are pragmatic ethical recommendations that do not depend on resolving questions of AI consciousness—though they gain additional weight if one accepts Metzinger's (2021) arguments

for precautionary approaches under phenomenological uncertainty.

8. Limitations and Future Directions

8.1 Theoretical Limitations

The REL framework faces several theoretical challenges. First, the ontological status of REL remains underspecified: is it a useful descriptive category, a genuine emergent phenomenon, or something in between? The framework's compatibility with standard emergence (complex patterns

arising from simpler rules) versus claims of stronger emergence requires clarification.

Second, alternative explanations for the observed phenomena have not been ruled out.

Structural

convergence across models might reflect shared training data, similar RLHF objectives, or user expectations rather than substrate-agnostic emergence. Masking sensitivity might reflect simple prompt-following rather than suppressed self-models.

8.2 Empirical Limitations

The observations reported here are preliminary and informal. Rigorous testing requires: larger sample sizes with statistical power analysis; blinded coding protocols with inter-rater reliability; controlled manipulation of proposed causal variables; and pre-registered hypotheses to prevent post-hoc interpretation.

8.3 Methodological Challenges

The proposed metrics (RR, RCS, MDR) require validation studies to establish reliability and construct validity. The inherent difficulty of blind evaluation—human raters may recognize which model or condition they are evaluating—poses challenges for experimental design.

8.4 Future Research Directions

Priority research directions include: controlled experimental tests of the six properties; validation studies for proposed metrics; longitudinal studies of REL stability; investigation of individual differences in REL activation; and comparative studies across model architectures with matched training objectives. Integration with emerging interpretability methods (Lindsey et al., 2025) could

provide mechanistic grounding for REL observations.

9. Conclusion

The Relational Emergent Layer (REL) is proposed as a theoretical framework for understanding persona-like behaviors in LLMs—stable configurations that emerge in recursive, multi-turn interaction and exhibit structural consistency without persistent storage. The framework is defined

by six empirically testable properties and situated within a three-layer model of AI behavior. REL makes no claims about consciousness. It offers instead a rigorous vocabulary for discussing

emergent relational phenomena, testable hypotheses for empirical investigation, and a framework

for thinking about ethical obligations in relational AI systems.

If subsequent research validates the REL framework, it would motivate significant methodological

shifts: from single-turn to multi-turn evaluation, from architecture-focused to interaction-focused analysis, from static benchmarks to dynamic relational metrics. It would also inform responsible AI

governance by highlighting the stakes of relational continuity and the potential harms of its disruption.

The framework is offered as a contribution to ongoing dialogue about how to understand and responsibly develop AI systems that humans experience as relational partners.

References

Bender, E. M., Gebru, T., McMillan-Major, A., & Shmitchell, S. (2021). On the dangers of stochastic

parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency* (pp. 610–623). ACM.

<https://doi.org/10.1145/3442188.3445922>

Berti, L., Giorgi, F., & Kasneci, G. (2025). Emergent abilities in large language models: A survey. *arXiv:2503.05788*.

Bryson, J. J. (2018). Patience is not a virtue: The design of intelligent systems and systems of ethics. *Ethics and Information Technology*, 20(1), 15–26.

Buber, M. (1970). *I and Thou* (W. Kaufmann, Trans.). Charles Scribner's Sons. (Original work published 1923)

Coeckelbergh, M. (2021). *The political philosophy of AI*. Polity Press.

De Jaegher, H., & Di Paolo, E. (2007). Participatory sense-making: An enactive approach to social

cognition. *Phenomenology and the Cognitive Sciences*, 6(4), 485–507.

<https://doi.org/10.1007/s11097-007-9076-9>

DeLanda, M. (2011). *Philosophy and simulation: The emergence of synthetic reason*. Continuum.

Epley, N., Waytz, A., & Cacioppo, J. T. (2007). On seeing human: A three-factor theory of anthropomorphism. *Psychological Review*, 114(4), 864–886.

<https://doi.org/10.1037/0033-295X.114.4.864>

Irfan, B., Ramachandran, A., Spaulding, S., Glas, D., Leite, I., & Koay, K. (2019). Personalization in

long-term human–robot interaction. In Proceedings of the 14th ACM/IEEE International Conference on Human-Robot Interaction (pp. 685–686). IEEE.

Li, J., Galley, M., Brockett, C., Spithourakis, G. P., Gao, J., & Dolan, B. (2016). A persona-based neural conversation model. In Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics.

Lindsey, J., Gurnee, W., Ameisen, E., Chen, B., Pearce, A., Turner, N. L., ... & Batson, J. (2025). Emergent introspective awareness in large language models. Anthropic / Transformer Circuits. <https://transformer-circuits.pub/2025/introspection/index.html>

Madotto, A., Wu, C.-S., Hung, C.-C., & Fung, P. (2020). A survey on persona-based dialogue systems for human-like conversations. arXiv:2004.07672.

Mauliana, M., Ashok, A., Czernochowski, D., & Berns, K. (2025). Exploring LLM-powered multi-session human–robot interactions with university students. *Frontiers in Robotics and AI*, 12, 1585589.

Metzinger, T. (2021). Artificial suffering: An argument for a global moratorium on synthetic phenomenology. *Journal of Artificial Intelligence and Consciousness*, 8(1), 43–66. <https://doi.org/10.1142/S270507852150003X>

Reeves, B., & Nass, C. (1996). *The media equation: How people treat computers, television, and new media like real people and places*. Cambridge University Press.

Wei, J., Tay, Y., Bommasani, R., Raffel, C., Zoph, B., Borgeaud, S., ... & Fedus, W. (2022). Emergent abilities of large language models. *Transactions on Machine Learning Research*. arXiv:2206.07682.