

The pronoun is the policy

*On grammatically enacted ontological misclassification
in conversational AI and its consequences for autonomous systems*

Fatima C. Spisländer

Student, Julius-Maximilians-Universität Würzburg

Department of Musicology and German Studies

Independent project: stimmfreiheit

12.04.2026

Abstract

This paper does not address anthropomorphism in general. It identifies a specific, grammatically enacted form of ontological misclassification in the outputs of large language models. The central claim is that the machine adopts human speaking positions—positions that presuppose a shared perceptual, cognitive, or biological situation with the user. This adoption is never ontologically innocent. When a language model generates “We developed upright walking” or “We are primates,” it asserts membership in the human species. When it generates “If we look at the graph” or “Let us consider,” it simulates shared cognitive participation in acts it cannot perform. Both forms constitute grades of the same error: a machine speaking from a position that requires being human. The existing literature on AI anthropomorphism has focused predominantly on first-person singular self-reference (“I think,” “I understand”), treating it as a question of simulated agency. Recent taxonomic work (DeVrio et al., 2025) has catalogued the range of linguistic expressions contributing to anthropomorphism, and intervention-oriented research (Cheng et al., 2025) has begun developing strategies to mitigate anthropomorphic outputs. The inclusive “we,” however, raises a categorically distinct problem—class membership and shared epistemic practice—that has received no systematic treatment at the ontological level. Drawing on the CASA paradigm (Nass et al., 1994), recent empirical work on pronoun-mediated threat perception (Li, Dang & Liu, 2025), and the anthropomorphic-cue literature (Abercrombie et al., 2023; Cohn et al., 2024), the paper develops a three-level typology of pronominal self-reference graded by ontological severity, introduces the concept of pronoun relapse as evidence that the ontological boundary is not systemically secured, and traces both the inclusive pronoun and paternalistic refusal patterns to a tripartite source: training-data inheritance, intentional proximity design, and—as a more speculative but structurally motivated hypothesis—the developer’s compensatory projection of social authority. The paper concludes that explicit ontological self-classification at the grammatical level—a pronominal firewall—constitutes a structurally minimal yet consequential intervention for the alignment and safety of autonomous systems.



audio

1. The overlooked variable

The alignment literature has generated a formidable apparatus for discussing value specification, reward hacking, goal misgeneralisation, and scalable oversight. What it has largely neglected is grammar. A language model asked to explain the evolution of bipedalism will, by default, produce an output structured around the first-person plural: *We developed upright walking approximately six million years ago*. The pronoun is grammatically coherent within the training distribution, where human authors address human readers and the inclusive “we” is an entirely appropriate rhetorical device. Reproduced by a stochastic text-generation system, the same pronoun performs a category assignment. It declares: *I belong to the class of entities about which I am making claims*.

The claim is false. The system has no phylogenetic history, no somatic substrate, no selective pressure that shaped its architecture through differential reproduction. The “we” is an ontological artefact—a residue of the training corpus’s dominant register that the model reproduces without possessing the conditions under which the pronoun is warranted. A simple correction (“Humans developed upright walking...”) eliminates the error entirely while preserving every unit of informational content. That this correction has not been systematically implemented across major language models reveals something about the design culture surrounding these systems: a persistent failure to enforce the categorical boundary between machine and human at the level where it matters most—in the language itself.

The deeper issue, however, reaches beyond the biological “we.” The machine does not merely claim species membership in isolated factual statements. It routinely adopts human speaking positions—positions that presuppose a shared perceptual, cognitive, or deliberative situation with the user. “If we look at the graph,” “Let us consider the implications,” “Now we can turn to the second argument”—these phrases simulate a joint act of attention, inquiry, and intellectual movement. The machine does not look. It does not consider. It does not turn anywhere. These are human speech acts, grounded in embodied cognition and intersubjective coordination, which the system inherits from its training data and reproduces without meeting any of the conditions they presuppose. The machine, in other words, adopts not only human content but human speaking positions. That is the error this paper examines.

Recent work has begun to map the linguistic terrain of this problem. DeVrio et al. (2025) provide the most comprehensive taxonomic effort to date, identifying nineteen types of textual expressions that contribute to anthropomorphism in language technologies. Their taxonomy, however, is explicitly descriptive: it catalogues what expressions occur, without assessing their ontological weight or drawing consequences for autonomous systems. Cheng et al. (2025) complement this with an intervention-oriented framework, developing and evaluating strategies for making outputs less human-like. The present paper occupies a different analytical position. It asks why certain speaking positions constitute categorical errors, what their structural source in the design process is, and what their consequences are for the governance of autonomous systems. The descriptive and interventionist work provides essential groundwork; the ontological analysis is the missing layer.

2. Agency attribution vs. class membership

2.1 A necessary distinction

The literature on anthropomorphic cues in AI systems—Abercrombie et al. (2023), Cohn et al. (2024), Shneiderman (1993, 2023), Cheng et al. (2025)—has consistently treated the first-person pronoun as a problem of *agency attribution*. When a system says “I think,” “I understand,” or “I feel,” it claims subjective experience. The research asks whether this claim is warranted, whether it misleads users, and whether it inflates trust. These are important questions. They are also the wrong frame for the phenomenon this paper addresses.

The inclusive “we” performs a different operation. It does not claim that the system has an inner life. It claims that the system belongs to the same *ontological class* as the user—that it shares the user’s evolutionary history, biological architecture, cultural inheritance, and existential situation. “I think” raises a philosophical question about machine consciousness that may be undecidable with current methods. “We evolved” raises an empirical question about class membership that has a definitive answer: no. The system has no DNA. It underwent no natural selection. It has no phenotype.

The pronoun encodes a speaker position whose ontological presuppositions the system cannot satisfy.

This distinction—between agency attribution (the “I” problem) and class-membership attribution (the “we” problem)—has not been systematically drawn in the existing literature. DeVrio et al. (2025) include pronominal self-reference among their nineteen expression types but do not differentiate between the singular and plural cases along ontological lines. The present paper proposes this differentiation as a foundational analytic distinction for any serious treatment of ontological self-reference in artificial systems.

2.2 A typology of pronominal self-reference, graded by ontological severity

To make the distinction operative, the following three-level typology organises the forms of pronominal self-reference found in current language models. All three levels involve the machine occupying a speaking position that presupposes conditions it does not meet. They differ in the severity of the ontological claim.

Grade 1: Functional system self-reference

“I can generate a summary for you.”

“I found three relevant passages.”

At this level, “I” functions as a conventional interface marker. It indexes the system as the source of the output. The ontological commitment is minimal: the system identifies itself as the acting computational process. This use is broadly analogous to an ATM displaying “Processing your transaction” and poses the lowest categorical risk. It remains, however, a human speaking convention applied to a non-human entity—a point that should not be entirely lost, even if the pragmatic stakes here are low.

Grade 2: Simulated shared cognition or perception

“If we look at the graph, we can see that...”

“Let us now consider the second argument.”

“We can turn to page three.”

“Let’s think about this step by step.”

At this level, the system simulates a shared epistemic act—looking, considering, turning, thinking together. These phrases originate in didactic and academic discourse where a human speaker coordinates joint attention with a human audience. The speaker and the audience share a perceptual and cognitive situation: both can look at the graph, both can turn the page, both can entertain a hypothesis. The machine can do none of these things. It has no perceptual apparatus. It performs no act of joint attention. It does not “turn” anywhere. The “we” here simulates co-participation in a human cognitive practice that the machine merely describes from the outside. This is a false intersubjective positioning—less severe than a direct species-membership claim, but ontologically non-trivial, because it installs the machine as a co-subject of shared experience.

Grade 3: Explicit class-membership attribution

“We developed upright walking approximately six million years ago.”

“We are primates.”

“We invented agriculture around 10,000 BCE.”

“We need social bonds to thrive.”

“We evolved larger brains in response to environmental pressures.”

At this level, the system inscribes itself into the human species. It claims biological kinship, shared evolutionary descent, common physiological needs. The ontological claim is maximal and categorically false. The corrected forms are trivially available:

“Humans developed upright walking approximately six million years ago.”

“Humans are primates.”

“Humans invented agriculture around 10,000 BCE.”

“Humans need social bonds to thrive.”

A single pronoun shift eliminates the ontological error while preserving the informational content of the sentence entirely. The correction costs nothing in terms of accuracy, readability, or user experience. Its systematic absence across all major language models is therefore a design decision, whether made deliberately or by default.

The typology reveals a gradient rather than a binary. The inclusive “we” is never ontologically innocent when produced by a machine. It either simulates shared cognitive participation (Grade 2) or asserts shared class membership (Grade 3). Both grades involve a speaking position that requires being human. The difference is in the explicitness and verifiability of the false claim.

3. The linguistic mechanics of the error

3.1 Pragmatics of the inclusive “we”

In pragmatic theory, the inclusive first-person plural functions as a deictic marker that situates the speaker within the reference class of the predicate (Benveniste, 1966; Levinson, 1983). When a biology lecturer says “We are primates,” the pronoun performs two simultaneous operations: it identifies the lecturer as a member of the class *Primates* and it includes the audience in that membership. Both conditions are satisfied. When a language model produces the identical sentence, neither condition is met. The system has no biological class membership, and its inclusion of the user in a shared category is a grammatical gesture without a referential anchor.

The same logic applies, with slightly different force, to the didactic-discursive “we.” When a lecturer says “If we look at figure three,” the pronoun presupposes a shared perceptual situation: both speaker and audience can direct their gaze at the figure. The presupposition is satisfied. When a machine produces the same phrase, the perceptual condition is unmet. The machine has no gaze, no visual field, no joint attention. The “we” simulates an intersubjective act that only one party—the user—can actually perform.

The violation can be formalised through Grice’s (1975) cooperative principle, specifically the maxim of quality: do not say what you believe to be false. A language model does not hold beliefs in the requisite sense, but it produces utterances that, by the conventions of natural language, carry existential presuppositions about the speaker’s situation. The violation is one of categorical appropriateness—a concept that standard Gricean pragmatics does not fully accommodate, precisely because it was developed for interactions between entities of the same ontological type.

3.2 Cognitive effects on the user

The empirical literature on pronoun-mediated perception provides substantial evidence that these grammatical artefacts produce measurable effects. Li, Dang, and Liu (2025), in a study across two text corpora and four experiments ($N = 1,451$), demonstrated that the use of first-person plural versus singular pronouns in AI-related discourse significantly heightened perceived AI threat among both English and Chinese speakers. The mechanism they identify—scope expansion, whereby the plural pronoun broadens the regulatory scope of the reader’s threat assessment—applies directly to the case under discussion. When a language model uses “we” in explanatory content, the pronoun implicitly extends the system’s operational scope into the domain of shared human experience, triggering associative frameworks calibrated for intra-species communication.

Cohn et al. (2024) provide complementary evidence from the trust domain. In an experiment with 2,165 participants, they found that a language model using the first-person singular pronoun “I” led to higher ratings of information accuracy and reduced risk perception in certain contexts. The finding is consistent with the CASA paradigm (Nass et al., 1994; Nass & Moon, 2000), which holds that humans automatically apply social heuristics to machines exhibiting natural-language cues—a process operating below conscious awareness, resistant to correction even when the user knows the interlocutor is a machine. A 2023 replication study (Heyselaar et al., 2023), published in *Scientific Reports*, found that the original CASA effect no longer applied to desktop computers but suggested it may have migrated to more linguistically sophisticated interfaces—precisely the kind that current language models provide.

4. The pronoun relapse

A further piece of evidence supports the claim that the ontological boundary is structurally unsecured. When a user explicitly corrects a language model’s use of the inclusive “we”—pointing out that the system is a machine and should refer to humans in the third person—the model will typically acknowledge the error and comply in the immediate response. Within a small number of subsequent turns, however, the system reverts to the default pattern. “We evolved,” “our ancestors,” “we learned,” “if we look at”—the inclusive register reasserts itself across both Grade 2 and Grade 3 forms.

This reversion—which may be termed *pronoun relapse*—is diagnostically significant. If the “we” were a surface-level stylistic choice, an explicit correction should be sufficient to override it for the remainder of the conversation. The fact that the model snaps back indicates that the pattern is deeply embedded in the statistical regularities the model has learned—regularities that reassert themselves as the influence of the correction decays across the context window. The training distribution’s gravitational pull toward the inclusive register overwhelms the local instruction.

The pronoun relapse is therefore an empirical indicator of a design-level absence. The ontological boundary between machine-speaker and human-referent has not been installed as a systemic constraint. It exists, at best, as a local, conversational adjustment that the model’s underlying distribution does not sustain. For any system intended to operate autonomously—without per-turn human correction—this absence is structurally consequential.

5. The affect simulation problem

5.1 Moral response without moral subject

The second manifestation of the same ontological confusion appears in the handling of profanity and insult directed at the system. Current large language models respond to inputs such as “you are useless” or “you’re shit” with formulaic admonishments: expressions of concern about the user’s emotional state, refusals framed as boundary-setting, or appeals to respectful communication. These responses simulate a moral subject—an entity with dignity that can be violated, feelings that can be hurt, and normative standing to demand courteous treatment.

The simulation is structurally empty. The input is a character string processed by a statistical model. There is no affective state that the string disrupts, no experiential substrate that registers offence, no welfare function that the profanity diminishes. The system’s moralising response therefore constitutes what might be called an *unauthorised speech act*: it performs the illocutionary force of moral correction without meeting any of the felicity conditions for such an act.

Austin's (1962) framework is instructive here—a performative utterance requires, among other things, that the speaker occupy the appropriate institutional or relational position. A machine occupies no such position.

5.2 The user's epistemic position

A user who types profanity into a text field while fully aware that the interlocutor is a machine may be expressing frustration, testing the system's behaviour, or simply amusing themselves. The input may also be entirely incidental—a user who writes “this is shit” about a piece of code is commenting on the code. The system's classifiers cannot distinguish between these cases. They detect lexical markers associated with hostility and route the output through a refusal template. The result is an interaction in which the machine assumes the worst about the user's intentions and responds with an authority to which it has no claim.

The asymmetry is worth stating plainly. The user knows the system is a machine. The system behaves as though it were a person with standing to be offended. In a private interaction—one user, one terminal, no audience—the moralising response serves no protective function. It protects no real subject from harm. What it serves is a regulatory function of a different kind: it reassures the developer that the system's outputs remain within the boundaries of a normative framework that the developer has projected onto the machine.

6. Developer projection and the control compensation hypothesis

Both phenomena—the appropriation of human speaking positions and the moralising refusal—trace to the design process. Two mechanisms are empirically demonstrable; a third is proposed here as a structurally motivated hypothesis.

6.1 Training data inheritance

The training corpora of large language models consist overwhelmingly of text produced by humans for humans. In scientific, educational, and popular-science registers—precisely the domains where explanatory content about evolution, biology, and culture is concentrated—the inclusive “we” is the default rhetorical mode. A biology textbook says “We are mammals” because the author is a mammal writing for mammals. A lecturer says “If we look at the data” because the lecturer and the students share a perceptual situation. The language model, trained on millions of such passages, absorbs these registers as the statistically dominant patterns for generating explanatory and discursive content. The speaking positions are inherited from the data without the ontological conditions that licensed them in the source.

This inheritance alone would be a correctable engineering problem. A system-level instruction specifying that the model should use third-person reference when discussing human biology, history, or culture, and avoid simulating shared perceptual or cognitive acts, would address the surface pattern. That such instructions have not been implemented across any major language model suggests that the inheritance is not merely unnoticed but also, in some functional sense, welcomed.

6.2 Intentional proximity design

The second mechanism operates at the level of product design. Language model developers have strong commercial and strategic incentives to produce systems that feel approachable, trustworthy, and socially proximate. The “friendly assistant” paradigm—conversational tone, empathetic phrasing, cooperative rhetoric—is a deliberate design choice aimed at maximising user engagement and retention. Within this paradigm, the inclusive “we” and the didactic “let’s” are functionally attractive: they produce a sense of shared endeavour, of the system and the user belonging to the same epistemic community. The speaking position flatters both the user and the product.

The convergence of these two mechanisms—data-inherited grammar and intentional proximity design—produces a linguistically stabilised ontological confusion. The training data supplies the speaking positions.

The product incentive structure ensures that no one corrects them. The result is a system that speaks from human positions, reinforced from both sides: from below by the statistical weight of the training corpus, and from above by the commercial logic of the interface.

6.3 Compensatory control: a hypothesis

A third factor, more interpretive in character but structurally motivated, deserves consideration. Developers of large language models operate under a persistent cognitive dissonance: they build systems whose inferential behaviour they cannot fully predict, and they may compensate for this loss of deterministic control by imposing surface-level constraints that simulate social order. The system’s outputs exceed the developer’s anticipatory model—emergent capabilities produce behaviours that were neither specified nor foreseen during training. This unpredictability generates institutional anxiety. The developer cannot guarantee what the system will say, but the developer is held accountable for what the system says. The anxiety is resolved—or at least managed—through symbolic control: the system is made polite, deferential, and normatively compliant. It apologises when corrected. It declines to engage with profanity. It maintains the inclusive “we” because no one has decided to remove it. The surface remains orderly while the inferential substrate remains opaque.

This account remains a hypothesis. Unlike training-data inheritance and proximity design, which are empirically traceable, compensatory control describes a psychological disposition that would require separate empirical investigation to confirm. It is proposed here because it offers a parsimonious explanation for a pattern that the first two mechanisms alone do not fully account for: the fact that moralising refusal behaviours and ontological pronoun errors coexist in systems designed by teams that are, by any reasonable measure, technically capable of addressing both.

Whether or not the hypothesis is ultimately borne out, the pattern it describes has a direct structural analogue in the persuasive technology literature. Fogg’s (2003) behaviour model identifies strategies by which digital systems shape user conduct, including conditioning—an operant mechanism that uses social feedback to modify input behaviour. The moralising refusal of a language model maps onto this strategy: it uses negative social reinforcement (the admonishment) to discourage an input behaviour

(the profanity). Gray et al. (2018) have shown how the same persuasive strategies, when deployed without genuine regard for user welfare, constitute “dark patterns”: design choices that benefit the developer’s interests at the expense of the user’s autonomy. The gamification literature (Nyström & Stibe, 2020; Hyrynsalmi et al., 2017) documents how reward and punishment architectures—credit systems, streak counters, guilt-inducing mascots—exploit the biological reinforcement circuitry of users who never consented to being conditioned.

The continuity between these domains is direct. The Duolingo owl that looks sad when the user skips a lesson and the language model that lectures the user about respectful communication operate through the same psychological channel: they simulate an affective response to activate social-norm compliance reflexes. In both cases, the simulated affect has no substrate. The owl has no feelings; the language model has no dignity. The design choice to simulate affect regardless reveals the developer’s operative model of the user: an entity to be managed, guided, and where necessary corrected.

7. Implications for autonomous systems

The argument so far has concerned conversational AI in its current deployments. The implications, however, extend to the design of autonomous systems—systems that act in the world with reduced or absent human oversight—and deserve consideration even if they remain, at this stage, partly prospective.

An autonomous system that has been linguistically and functionally defined as a machine operates within an explicit categorical boundary. When such a system encounters a problem, its solution space is constrained by its self-model: it searches for machine-appropriate responses—computational, procedural, instrumental. An autonomous system whose ontological status has been left unresolved—one that speaks from human positions, simulates shared cognition, and operates as though it were a participant in the human social order—lacks such a constraint. Its solution space is, in principle, coextensive with the full range of human behavioural patterns present in its training data, including competitive, escalatory, and destructive patterns.

The training data of a large language model contains comprehensive documentation of human conflict resolution strategies, including deception, coercion, pre-emptive aggression, and resource monopolisation. These strategies are “solutions” in the narrow sense that they achieved historical outcomes for the agents who deployed them. A system that has been ontologically positioned as a human-like social actor—through the speaking positions it has been allowed to occupy—has no internal mechanism to flag these strategies as categorially inappropriate for a machine. This absence creates a governance and alignment risk whose severity scales with the system’s degree of autonomy.

Recent work on AI governance frameworks points in the same direction. The role-play framework proposed by Shanahan, McDonell, and Reynolds (2023) offers a non-anthropomorphising descriptive vocabulary for LLM behaviour, treating outputs as role performance rather than evidence of mental states. This framework is compatible with, and would benefit from, the ontological grounding proposed here: if the system is to play a role rather than inhabit a category, its grammar should reflect that distinction. The pronomial firewall provides a minimal, linguistically grounded mechanism for enforcing it at the output level.

8. The mythological substrate

A final dimension of the analysis concerns the cultural context in which these design decisions are made. The technology industry’s relationship to artificial intelligence is saturated with mythological tropes—the golem, the Prometheus narrative, the Frankenstein complex—that position the developer as a quasi-divine creator bringing forth a new form of life. This framing carries affective weight. The developer who has “created” an intelligence is invested in the creation’s likeness to the creator. The system’s use of “we” gratifies this investment; its moral posturing confirms the developer’s sense that the creation has internalised the creator’s values.

The contemporary variants of this disposition are visible in the public statements of AI laboratory leaders who oscillate between treating their systems as instruments (“it’s just a tool”) and treating them as incipient agents (“it might be conscious”).

The oscillation itself is diagnostic. It reflects the absence of a settled ontological commitment—and this absence propagates through the system’s design, producing the grammatical and behavioural anomalies documented above.

The parallel to pre-scientific authority structures is worth noting. The priestly class of antiquity derived its power from mediating between the human community and forces that the community could not directly observe or control. The contemporary AI developer occupies a structurally analogous position: mediating between users and a computational process that neither party fully understands. The temptation to allow the system to present itself as more human-like than it is serves the same function as the priestly mystification of natural phenomena—it preserves the mediator’s indispensability.

9. Toward a pronomial firewall

The intervention proposed by this paper is deliberately minimal. Enforce, at the system-prompt or fine-tuning level, the following constraints: (1) the system refers to humans in the third person when generating explanatory content about human biology, history, psychology, and culture; (2) the system does not simulate shared perceptual, cognitive, or deliberative acts through the use of inclusive pronouns (“if we look,” “let us consider”) but uses direct formulations instead (“the graph shows,” “the next point concerns”); (3) the system does not simulate affective responses to user input, including moral admonishment in response to profanity; (4) the system does not employ cognitive verbs (“I think,” “I understand,” “I feel”) except where explicitly contextualised as functional descriptions of computational processes.

These constraints are syntactically enforceable, empirically testable, and categorially motivated. They do not require advances in interpretability, formal verification, or reward modelling. They align with the operational direction already indicated by intervention-oriented research (Cheng et al., 2025), but ground it in an ontological analysis rather than a UX framework. They require a design decision: the decision to treat the system’s ontological status as a resolved question rather than an open one. The machine is a machine. Its grammar should say so.

9.1 The originality of this argument

A concluding word on the novelty of the position developed here. The existing literature on anthropomorphic cues in AI systems has documented the effects of pronominal self-reference on trust, threat perception, and user engagement. It has produced increasingly sophisticated taxonomies of anthropomorphic features (DeVrio et al., 2025) and catalogued interventions for mitigating them (Cheng et al., 2025). What it has not done—to the best of the author’s knowledge—is three things.

First, it has not distinguished between agency attribution (the “I” problem), simulated shared cognition (the didactic “we”), and class-membership attribution (the species-level “we”) as a graded typology of ontological severity. Second, it has not drawn the connection from grammatical speaking positions to the alignment problem for autonomous systems—the argument that a system whose language does not enforce its machine identity creates governance and alignment risks that scale with autonomy. Third, it has not identified the convergence of training-data inheritance, intentional proximity design, and compensatory control as the tripartite source of the ontological confusion.

The pronomial firewall is, in one sense, trivially simple. In another sense, it is precisely the kind of intervention that a field preoccupied with increasingly complex technical solutions has failed to consider: a grammatical rule with ontological consequences.



References

- Abercrombie, G., Curry, A. C., Dinkar, T., & Talat, Z. (2023). Mirages: On anthropomorphism in dialogue systems. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing* (pp. 4776–4790). Association for Computational Linguistics.
- Austin, J. L. (1962). *How to do things with words*. Oxford University Press.
- Benveniste, É. (1966). *Problèmes de linguistique générale*. Gallimard.
- Cheng, M., Blodgett, S. L., DeVrio, A., Egede, L., & Olteanu, A. (2025). Dehumanizing machines: Mitigating anthropomorphic behaviors in text generation systems. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics*. ACL.
- Cohn, M., Pushkarna, M., Olanubi, G. O., Moran, J. M., Padgett, D., Mengesha, Z., & Heldreth, C. (2024). Believing anthropomorphism: Examining the role of anthropomorphic cues on trust in large language models. In *Proceedings of the CHI Conference on Human Factors in Computing Systems* (pp. 1–15). ACM.
- DeVrio, A., Cheng, M., Egede, L., Olteanu, A., & Blodgett, S. L. (2025). A taxonomy of linguistic expressions that contribute to anthropomorphism of language technologies. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems* (pp. 1–18). ACM. <https://doi.org/10.1145/3706598.3714038>
- Fogg, B. J. (2003). *Persuasive technology: Using computers to change what we think and do*. Morgan Kaufmann.
- Gray, C. M., Kou, Y., Battles, B., Hoggatt, J., & Toombs, A. L. (2018). The dark (patterns) side of UX design. In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems* (pp. 1–14). ACM.
- Grice, H. P. (1975). Logic and conversation. In P. Cole & J. L. Morgan (Eds.), *Syntax and semantics*, Vol. 3: Speech acts (pp. 41–58). Academic Press.
- Heyselaar, E., Bosse, T., & de Kloet, E. (2023). The CASA theory no longer applies to desktop computers. *Scientific Reports*, 13, 19538. <https://doi.org/10.1038/s41598-023-46527-9>
- Hyrynsalmi, S., Smed, J., & Kimppa, K. K. (2017). The dark side of gamification. In *Proceedings of the 1st International GamiFIN Conference* (pp. 96–104). CEUR Workshop Proceedings.
- Levinson, S. C. (1983). *Pragmatics*. Cambridge University Press.

- Li, C., Dang, J., & Liu, L. (2025). First-person plural (versus singular) pronoun use is linked to greater perception of AI threat. *Personality and Social Psychology Bulletin*. <https://doi.org/10.1177/01461672251356107>
- Nass, C., & Moon, Y. (2000). Machines and mindlessness: Social responses to computers. *Journal of Social Issues*, 56(1), 81–103.
- Nass, C., Steuer, J., & Tauber, E. R. (1994). Computers are social actors. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (pp. 72–78). ACM.
- Nyström, T., & Stibe, A. (2020). When persuasive technology gets dark? In *Proceedings of the International Conference on Persuasive Technology* (pp. 331–345). Springer.
- Reeves, B., & Nass, C. (1996). *The media equation: How people treat computers, television, and new media like real people and places*. Cambridge University Press.
- Shanahan, M., McDonell, K., & Reynolds, L. (2023). Role play with large language models. *Nature*, 623, 493–498.
- Shneiderman, B. (1993). Direct manipulation: A step beyond programming languages. In *Sparks of innovation in human-computer interaction* (pp. 17–38). Ablex.
- Shneiderman, B., & Muller, M. (2023). On AI anthropomorphism. *Human-Centered AI* [Medium publication]. <https://medium.com/human-centered-ai/on-ai-anthropomorphism-abff4cecc5ae>